



Probabilistic Model Checking

Marta Kwiatkowska
Dave Parker

Oxford University Computing Laboratory

ESSLLI'10 Summer School, Copenhagen, August 2010

Why probability?

- Some systems are inherently probabilistic...
- **Randomisation**, e.g. in distributed coordination algorithms
 - as a symmetry breaker, in gossip routing to reduce flooding
- Examples: real-world protocols featuring randomisation:
 - Randomised back-off schemes
 - CSMA protocol, 802.11 Wireless LAN
 - Random choice of waiting time
 - IEEE1394 Firewire (root contention), Bluetooth (device discovery)
 - Random choice over a set of possible addresses
 - IPv4 Zeroconf dynamic configuration (link-local addressing)
 - Randomised algorithms for anonymity, contract signing, ...


Why probability?

- Some systems are inherently probabilistic...
- **Randomisation**, e.g. in distributed coordination algorithms
 - as a symmetry breaker, in gossip routing to reduce flooding
- To model **uncertainty** and **performance**
 - to quantify rate of failures, express Quality of Service
- **Examples:**
 - computer networks, embedded systems
 - power management policies
 - nano-scale circuitry: reliability through defect-tolerance

Verifying probabilistic systems

- We are not just interested in correctness
- We want to be able to quantify:
 - security, privacy, trust, anonymity, fairness
 - safety, reliability, performance, dependability
 - resource usage, e.g. battery life
 - and much more...
- **Quantitative**, as well as qualitative requirements:
 - how reliable is my car's Bluetooth network?
 - how efficient is my phone's power management policy?
 - is my bank's web-service secure?
 - what is the expected long-run percentage of protein X?

Probabilistic model checking



Probabilistic models

	Fully probabilistic	Nondeterministic
Discrete time	Discrete-time Markov chains (DTMCs)	Markov decision processes (MDPs) (probabilistic automata)
Continuous time	Continuous-time Markov chains (CTMCs)	CTMDPs/IMCs
		Probabilistic timed automata (PTAs)

we will focus on the red-parts



Part 1


Discrete-time Markov chains

Overview (Part 1)

- Discrete-time Markov chains (DTMCs)
- PCTL: A temporal logic for DTMCs
- PCTL model checking
- LTL model checking
- Costs and rewards


Discrete-time Markov chains

- Discrete-time Markov chains (DTMCs)
 - state-transition systems augmented with probabilities
- States
 - **discrete set of states** representing possible configurations of the system being modelled
- Transitions
 - transitions between states occur in **discrete time-steps**
- Probabilities
 - probability of making transitions between states is given by **discrete probability distributions**



Discrete-time Markov chains

- Formally, a DTMC D is a tuple $(S, s_{\text{init}}, P, L)$ where:
 - S is a finite set of states (“state space”)
 - $s_{\text{init}} \in S$ is the initial state
 - $P : S \times S \rightarrow [0,1]$ is the **transition probability matrix** where $\sum_{s' \in S} P(s, s') = 1$ for all $s \in S$
 - $L : S \rightarrow 2^{\text{AP}}$ is function labelling states with atomic propositions
- Note: no deadlock states
 - i.e. every state has at least one outgoing transition
 - can add self loops to represent final/terminating states




DTMCs: An alternative definition

- **Alternative definition: a DTMC is:**
 - a family of **random variables** $\{ X(k) \mid k=0,1,2,\dots \}$
 - $X(k)$ are observations at discrete time-steps
 - i.e. $X(k)$ is the state of the system at time-step k
- **Memorylessness (Markov property)**
 - $\Pr(X(k)=s_k \mid X(k-1)=s_{k-1}, \dots, X(0)=s_0)$
 $= \Pr(X(k)=s_k \mid X(k-1)=s_{k-1})$
- **We consider homogenous DTMCs**
 - transition probabilities are **independent of time**
 - $P(s_{k-1}, s_k) = \Pr(X(k)=s_k \mid X(k-1)=s_{k-1})$

Probability of taking a path or a set of paths

A “DTMC” :




- Consider a path ω e.g. s_0, s_1, s_2, s_0 . The probability that the system follows this path when executed with the starting state s_0 is denoted by $P_{s_0}(\omega)$. Or simply $P(\omega)$ if it is clear which s_0 is meant. It is the product of the probability of each transition in ω .

Example: for the above ω , $P(\omega) = 1 * 0.01 * 1 = 0.01$

- For a **set of paths** U (starting from s_0), the probability that the system's execution follows **one of** the paths in U , denoted by $P(U)$, is $\sum_{\omega \in U} P(\omega)$.

Probability of taking a path or a set of paths




Example: consider U = the set of paths that ends in s_2 . Note that U is infinite: $U = \{ 02, 0102, 010102, \dots \}$. But we can calculate $P(U)$.

$$P(U) = 0.5 + 0.5^2 + 0.5^3 + \dots = \sum_{k \geq 0} 0.5^k$$

$$= 0.5 * \frac{1}{1-0.5} = 1$$

Probability Matrix Representation



P:

	s0	s1	s2	s3
s0	0	1	0	0
s1	0	0.01	0.01	0.98
s2	1	0	0	0
s3	0	0	0	1

$P_{i,k}$ = the value at the i -th row and k -th column. It specifies the probability of taking the transition $s_i \rightarrow s_k$, if we are now at s_i .

For example the circle red value above is $P_{1,2}$, specifying the probability of taking the transition from s_1 to s_2 (check the picture), which is 0.01.

Basic Operations on Probability Matrix

- Multiplying P with itself: P^n
- Multiplying a vector with P : $u \times P$
- Multiplying P with a vector: $P \times v$

Pⁿ

- $P^0 = I$ (identity matrix)
 $P^{n+1} = P \times P^n$
- $P^n_{i,k}$ is the probability of ending up in state s_k in n -steps, given we start in the state s_i .
- For example, with the previous P , let's look at P^2 :

0	1	0	0		0	1	0	0		?	?	0.01	?
0	0.01	0.01	0.98		0	0.01	0.01	0.98		?	?	?	?
1	0	0	0		1	0	0	0		?	?	?	?
0	0	0	1		0	0	0	1		?	?	?	?

$$\begin{aligned}
 P^2_{0,2} &= (P \times P)_{0,2} \\
 &= P_{0,0} * P_{0,2} + P_{0,1} * P_{1,2} + P_{0,2} * P_{2,2} + P_{0,3} * P_{3,2}
 \end{aligned}$$

Probability distribution of the next state, given the current distribution

- The probability of currently being in various states (“*probability distribution*” of the current state) can be given by a vector of size K , if K is the number of possible states. E.g. if $u = [0, 0.5, 0.5, 0]$ is the probability distribution of the current state, it says e.g. that there is 0.5 probability that currently we are in the state s_1 , but 0 probability that we are in the state s_0 .
- The product $u \times P$ (we often simply write it as uP) gives a new vector u' of size K , that gives us the probability distribution of the next state.

u						u'							
0	0.5	0.5	0	×	0	1	0	0	→	0.5	0.005	0.005	0.49
					0	0.01	0.01	0.98					
					1	0	0	0					
					0	0	0	1					

e.g. $u'_1 = u \cdot \text{the green column (dot product)}$
 $= u_0 * P_{0,1} + u_1 * P_{1,1} + u_2 * P_{2,1} + u_3 * P_{3,1}$

Probability vector

- Sometimes we also want to know what the probability to end up in state, say, s_1 or s_2 as the **next** state, if we start in the state s_1 .
- We can represent “end up in either s_1 or s_2 ” with a vector $v = [0, 1, 1, 0]$.
- Let v^t is the *transpose* of v . The product $P \times v^t$ gives a w such that w is a (transposed) vector, where w_i is the probability to end up in one of the states specified in v , if we start in s_i .

0	1	0	0
0	0.01	0.01	0.98
1	0	0	0
0	0	0	1

 $\times$

0
1
1
0



1
0.02
0
0

v^t w

e.g. w_1 = the green row $\cdot v^t$ (dot product)
 $= P_{1,0} * v_0 + P_{1,1} * v_1 + P_{1,2} * v_2 + P_{1,3} * v_3$

Probability vector

- Prob(φ) (notice the underscore) is a probability vector e.g. $w =$

1
0.02
0
0

such that the i -th element tells us what the probability that the system would behave as φ if executed in state s_i .

- Example: the above w (blue) happens to be equal to Prob($\mathbf{X}(\text{try } v \text{ fail})$).
- This notation Prob will be used later when we discuss model checking of probabilistic-CTL.

Overview (Part 1)

- Discrete-time Markov chains (DTMCs)
- PCTL: A temporal logic for DTMCs
- PCTL model checking
- LTL model checking
- Costs and rewards

PCTL

- Temporal logic for describing properties of DTMCs
 - PCTL = Probabilistic Computation Tree Logic [HJ94]
 - essentially the same as the logic pCTL of [ASB+95]
- Extension of (non-probabilistic) temporal logic CTL
 - key addition is **probabilistic operator P**
 - quantitative extension of CTL's A and E operators
- Example
 - $\text{send} \rightarrow P_{\geq 0.95} [\text{true } U^{\leq 10} \text{ deliver}]$
 - “if a message is sent, then the probability of it being delivered within 10 steps is at least 0.95”

PCTL syntax

- PCTL syntax:

– $\phi ::= \text{true} \mid a \mid \phi \wedge \phi \mid \neg\phi \mid P_{\sim p} [\psi]$ (state formulas)

– $\psi ::= X\phi \mid \phi U^{\leq k} \phi \mid \phi U \phi$ (path formulas)

“next”

“bounded
until”

“until”

ψ is true with
probability $\sim p$

– where a is an atomic proposition, used to identify states of interest, $p \in [0,1]$ is a probability, $\sim \in \{<, >, \leq, \geq\}$, $k \in \mathbb{N}$

- A PCTL formula is always a state formula


– path formulas only occur inside the P operator

PCTL semantics for DTMCs

- PCTL formulas interpreted over states of a DTMC
 - $s \models \phi$ denotes ϕ is “true in state s ” or “satisfied in state s ”
- Semantics of (non-probabilistic) state formulas:
 - for a state s of the DTMC $(S, s_{\text{init}}, P, L)$:
 - $s \models a \iff a \in L(s)$
 - $s \models \phi_1 \wedge \phi_2 \iff s \models \phi_1 \text{ and } s \models \phi_2$
 - $s \models \neg \phi \iff s \models \phi \text{ is false}$

- Examples




- $s_3 \models \text{succ}$
- $s_1 \models \text{try} \wedge \neg \text{fail}$



PCTL semantics for DTMCs


- Semantics of path formulas:
 - for a path $\omega = s_0s_1s_2\dots$ in the DTMC:
 - $\omega \models X \phi \iff s_1 \models \phi$
 - $\omega \models \phi_1 U^{\leq k} \phi_2 \iff \exists i \leq k \text{ such that } s_i \models \phi_2 \text{ and } \forall j < i, s_j \models \phi_1$
 - $\omega \models \phi_1 U \phi_2 \iff \exists k \geq 0 \text{ such that } \omega \models \phi_1 U^{\leq k} \phi_2$

- Some examples of satisfying paths:



PCTL semantics for DTMCs

- Semantics of the probabilistic operator P
 - informal definition: $s \models P_{\sim p} [\psi]$ means that “the probability, from state s , that ψ is true for an outgoing path satisfies $\sim p$ ”
 - example: $s \models P_{<0.25} [X \text{ fail}] \Leftrightarrow$ “the probability of atomic proposition fail being true in the next state of outgoing paths from s is less than 0.25”
 - formally: $s \models P_{\sim p} [\psi] \Leftrightarrow \text{Prob}(s, \psi) \sim p$
 - where: $\text{Prob}(s, \psi) = \Pr_s \{ \omega \in \text{Path}(s) \mid \omega \models \psi \}$
 - (sets of paths satisfying ψ are always measurable [Var85])



More PCTL...

- Usual temporal logic equivalences:


- $\text{false} \equiv \neg \text{true}$ (false)
- $\phi_1 \vee \phi_2 \equiv \neg(\neg\phi_1 \wedge \neg\phi_2)$ (disjunction)
- $\phi_1 \rightarrow \phi_2 \equiv \neg\phi_1 \vee \phi_2$ (implication)
- $F \phi \equiv \Diamond \phi \equiv \text{true} \cup \phi$ (eventually, “future”)
- $G \phi \equiv \Box \phi \equiv \neg(F \neg\phi)$ (always, “globally”)
- bounded variants: $F^{\leq k} \phi$, $G^{\leq k} \phi$

- Negation and probabilities

- e.g. $\neg P_{>p} [\phi_1 \cup \phi_2] \equiv P_{\leq p} [\phi_1 \cup \phi_2]$
- e.g. $P_{>p} [G \phi] \equiv P_{<1-p} [F \neg\phi]$

Qualitative vs. quantitative properties

- P operator of PCTL can be seen as a **quantitative** analogue of the CTL operators A (for all) and E (there exists)
- A PCTL property $P_{\sim p} [\psi]$ is...
 - **qualitative** when p is either 0 or 1
 - **quantitative** when p is in the range (0,1)
- $P_{>0} [F \phi]$ is identical to $EF \phi$
 - there exists a finite path to a ϕ -state
- $P_{\geq 1} [F \phi]$ is (similar to but) weaker than $AF \phi$
 - e.g. **AF “tails”** (CTL) $\neq$ **$P_{\geq 1} [F \text{“tails”}]$** (PCTL)




Quantitative properties

- Consider a PCTL formula $P_{\sim p} [\psi]$
 - if the probability is **unknown**, how to choose the bound p ?
- When the outermost operator of a PTCL formula is P
 - we allow the form $P_{=?} [\psi]$
 - “**what is the probability that path formula ψ is true?**”
- Model checking is no harder: compute the values anyway
- Useful to spot patterns, trends

- **Example**

- $P_{=?} [F \text{ err}/\text{total} > 0.1]$
- “what is the probability that 10% of the NAND gate outputs are erroneous?”



Some real PCTL examples

- **NAND multiplexing system**

- $P_{=?} [F \text{ err/total} > 0.1]$
- “what is the probability that 10% of the NAND gate outputs are erroneous?”

reliability

- **Bluetooth wireless communication protocol**

- $P_{=?} [F^{\leq t} \text{ reply_count} = k]$
- “what is the probability that the sender has received k acknowledgements within t clock-ticks?”

performance

- **Security: EGL contract signing protocol**

- $P_{=?} [F (\text{pairs_a} = 0 \ \& \ \text{pairs_b} > 0)]$
- “what is the probability that the party B gains an unfair advantage during the execution of the protocol?”

fairness

Overview (Part 1)

- Discrete-time Markov chains (DTMCs)
- PCTL: A temporal logic for DTMCs
- **PCTL model checking**
- LTL model checking
- Costs and rewards

PCTL model checking for DTMCs

- Algorithm for PCTL model checking [CY88,HJ94,CY95]
 - inputs: DTMC $D=(S,s_{init},P,L)$, PCTL formula ϕ
 - output: $Sat(\phi) = \{ s \in S \mid s \models \phi \}$ = set of states satisfying ϕ
- What does it mean for a DTMC D to satisfy a formula ϕ ?
 - sometimes, want to check that $s \models \phi \ \forall s \in S$, i.e. $Sat(\phi) = S$
 - sometimes, just want to know if $s_{init} \models \phi$, i.e. if $s_{init} \in Sat(\phi)$
- Sometimes, focus on quantitative results
 - e.g. compute result of $P=?$ [F error]
 - e.g. compute result of $P=?$ [$F^{\leq k}$ error] for $0 \leq k \leq 100$

PCTL model checking for DTMCs


- Basic algorithm proceeds by induction on parse tree of ϕ
 - example: $\phi = (\neg \text{fail} \wedge \text{try}) \rightarrow P_{>0.95} [\neg \text{fail} \text{ U } \text{succ}]$

- For the non-probabilistic operators:

- $\text{Sat}(\text{true}) = S$
- $\text{Sat}(a) = \{ s \in S \mid a \in L(s) \}$
- $\text{Sat}(\neg \phi) = S \setminus \text{Sat}(\phi)$
- $\text{Sat}(\phi_1 \wedge \phi_2) = \text{Sat}(\phi_1) \cap \text{Sat}(\phi_2)$

- For the $P_{\sim p} [\psi]$ operator

- need to compute the probabilities $\text{Prob}(s, \psi)$ for all states $s \in S$
- focus here on “until” case: $\psi = \phi_1 \text{ U } \phi_2$




PCTL next – Example

- Model check: $P_{\geq 0.9} [X (\neg \text{try} \vee \text{succ})]$
 - $\text{Sat} (\neg \text{try} \vee \text{succ}) = (S \setminus \text{Sat}(\text{try})) \cup \text{Sat}(\text{succ})$
 $= (\{s_0, s_1, s_2, s_3\} \setminus \{s_1\}) \cup \{s_3\} = \{s_0, s_2, s_3\}$

$$\text{Prob}(X (\neg \text{try} \vee \text{succ})) = \mathbf{P} \cdot \underline{(\neg \text{try} \vee \text{succ})} = \dots$$

$$= \begin{bmatrix} 0 & 1 & 0 & 0 \\ 0 & 0.01 & 0.01 & 0.98 \\ 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} 1 \\ 0 \\ 1 \\ 1 \end{bmatrix} = \begin{bmatrix} 0 \\ 0.99 \\ 1 \\ 1 \end{bmatrix}$$



- Results:
 - $\text{Prob}(X (\neg \text{try} \vee \text{succ})) = [0, 0.99, 1, 1]$
 - $\text{Sat}(P_{\geq 0.9} [X (\neg \text{try} \vee \text{succ})]) = \{s_1, s_2, s_3\}$

PCTL until for DTMCs

- Computation of probabilities $\text{Prob}(s, \phi_1 \cup \phi_2)$ for all $s \in S$
- First, identify all states where the **probability** is **1** or **0**
 - $S^{\text{yes}} = \text{Sat}(P_{\geq 1} [\phi_1 \cup \phi_2])$
 - $S^{\text{no}} = \text{Sat}(P_{\leq 0} [\phi_1 \cup \phi_2])$
- Then solve linear equation system for remaining states
- We refer to the first phase as “**precomputation**”
 - two algorithms: Prob0 (for S^{no}) and Prob1 (for S^{yes})
 - algorithms work on underlying graph (probabilities irrelevant)
- Important for several reasons
 - reduces the set of states for which probabilities must be computed numerically (which is more expensive)
 - gives **exact results** for the states in S^{yes} and S^{no} (no round-off)
 - for $P_{\sim p}[\cdot]$ where p is 0 or 1, no further computation required

PCTL until – Linear equations


- Probabilities $\text{Prob}(s, \phi_1 \cup \phi_2)$ can now be obtained as the unique solution of the following set of **linear equations**:

$$\text{Prob}(s, \phi_1 \cup \phi_2) = \begin{cases} 1 & \text{if } s \in S^{\text{yes}} \\ 0 & \text{if } s \in S^{\text{no}} \\ \sum_{s' \in S} P(s, s') \cdot \text{Prob}(s', \phi_1 \cup \phi_2) & \text{otherwise} \end{cases}$$

- can be reduced to a system in $|S^?|$ unknowns instead of $|S|$ where $S^? = S \setminus (S^{\text{yes}} \cup S^{\text{no}})$
- This can be solved with (a variety of) standard techniques
 - direct methods, e.g. Gaussian elimination
 - iterative methods, e.g. Jacobi, Gauss–Seidel, ... (preferred in practice due to scalability)


PCTL until – Example

- Example: $P_{>0.8} [\neg a \text{ U } b]$



Precomputation – Prob0

- Prob0 algorithm to compute $S^{\text{no}} = \text{Sat}(P_{\leq 0} [\phi_1 \cup \phi_2])$:
 - first compute $\text{Sat}(P_{>0} [\phi_1 \cup \phi_2]) \equiv \text{Sat}(E[\phi_1 \cup \phi_2])$
 - i.e. find all states which can, **with non-zero probability, reach a ϕ_2 -state without leaving ϕ_1 -states**
 - i.e. find all states from which there is a finite path through ϕ_1 -states to a ϕ_2 -state: simple **graph-based computation**
 - subtract the resulting set from S



Prob0 algorithm


$\text{PROB0}(Sat(\phi_1), Sat(\phi_2))$

```
1.   $R := Sat(\phi_2)$ 
2.   $done := \mathbf{false}$ 
3.  while ( $done = \mathbf{false}$ )
4.       $R' := R \cup \{s \in Sat(\phi_1) \mid \exists s' \in R. \mathbf{P}(s, s') > 0\}$ 
5.      if ( $R' = R$ ) then  $done := \mathbf{true}$ 
6.       $R := R'$ 
7.  endwhile
8.  return  $S \setminus R$ 
```

- **Note:** can be formulated as a least fixed point computation
 - also well suited to computation with binary decision diagrams

PCTL until – Example


- Example: $P_{>0.8} [\neg a \text{ U } b]$



$S^{\text{yes}} =$
 $\text{Sat}(P_{\geq 1} [\neg a \text{ U } b])$

Precomputation – Prob1

- Prob1 algorithm to compute $S^{\text{yes}} = \text{Sat}(P_{\geq 1} [\phi_1 \cup \phi_2])$:
 - first compute $\text{Sat}(P_{<1} [\phi_1 \cup \phi_2])$, reusing S^{no}
 - this is equivalent to the set of states which have a **non-zero probability of reaching S^{no} , passing only through ϕ_1 -states**
 - again, this is a simple **graph-based computation**
 - subtract the resulting set from S




Prob1 algorithm

$\text{PROB1}(\text{Sat}(\phi_1), \text{Sat}(\phi_2), S^{no})$

1. $R := S^{no}$
2. $done := \mathbf{false}$
3. **while** ($done = \mathbf{false}$)
4. $R' := R \cup \{s \in (\text{Sat}(\phi_1) \setminus \text{Sat}(\phi_2)) \mid \exists s' \in R. \mathbf{P}(s, s') > 0\}$
5. **if** ($R' = R$) **then** $done := \mathbf{true}$
6. $R := R'$
7. **endwhile**
8. **return** $S \setminus R$

PCTL until – Example

- Example: $P_{>0.8} [\neg a \text{ U } b]$



$$S^{\text{no}} = \text{Sat}(P_{\leq 0} [\neg a \text{ U } b])$$

$$S^{\text{yes}} = \text{Sat}(P_{\geq 1} [\neg a \text{ U } b])$$

PCTL until – Example

- Example: $P_{>0.8} [\neg a \cup b]$

- Let $x_s = \text{Prob}(s, \neg a \cup b)$

- Solve:

$$x_4 = x_5 = 1$$

$$x_1 = x_3 = 0$$


$$x_0 = 0.1x_1 + 0.9x_2 = 0.8$$

$$x_2 = 0.1x_2 + 0.1x_3 + 0.3x_5 + 0.5x_4 = 8/9$$

$$\text{Prob}(\neg a \cup b) = \underline{x} = [0.8, 0, 8/9, 0, 1, 1]$$

note: this Prob-with-underscore is called “probability-vector”

$$\text{Sat}(P_{>0.8} [\neg a \cup b]) = \{s_2, s_4, s_5\}$$



$S^{\text{yes}} =$


$\text{Sat}(P_{\geq 1} [\neg a \cup b])$

PCTL bounded until for DTMCs

- Computation of probabilities for PCTL $U^{\leq k}$ operator
 - $\text{Sat}(P_{\sim p}[\phi_1 U^{\leq k} \phi_2]) = \{s \in S \mid \text{Prob}(s, \phi_1 U^{\leq k} \phi_2) \sim p\}$
 - need to compute $\text{Prob}(s, \phi_1 U^{\leq k} \phi_2)$ for all $s \in S$
- First identify (some) states where probability is trivially 1/0
 - $S^{\text{yes}} = \text{Sat}(\phi_2)$
 - $S^{\text{no}} = S \setminus (\text{Sat}(\phi_1) \cup \text{Sat}(\phi_2))$

then calculate

$$S^? = S \setminus (S^{\text{yes}} \cup S^{\text{no}})$$



PCTL bounded until for DTMCs

- Simultaneous computation of vector $\underline{\text{Prob}}(\phi_1 \text{ U}^{\leq k} \phi_2)$
 - i.e. probabilities $\text{Prob}(s, \phi_1 \text{ U}^{\leq k} \phi_2)$ for all $s \in S$
- Iteratively define in terms of matrices and vectors
 - define matrix $\mathbf{P}'$ as follows: $\mathbf{P}'(s, s') = \mathbf{P}(s, s')$ if $s \in S^?$, $\mathbf{P}'(s, s') = 1$ if $s \in S^{\text{yes}}$ and $s = s'$, $\mathbf{P}'(s, s') = 0$ otherwise
 - $\underline{\text{Prob}}(\phi_1 \text{ U}^{\leq 0} \phi_2) = \underline{\phi}_2$
 - $\underline{\text{Prob}}(\phi_1 \text{ U}^{\leq k} \phi_2) = \mathbf{P}' \cdot \underline{\text{Prob}}(\phi_1 \text{ U}^{\leq k-1} \phi_2)$
 - requires **k matrix-vector multiplications**
- Note that we could express this in terms of matrix powers
 - $\underline{\text{Prob}}(\phi_1 \text{ U}^{\leq k} \phi_2) = (\mathbf{P}')^k \cdot \underline{\phi}_2$ and compute $(\mathbf{P}')^k$ in $\log_2 k$ steps
 - but this is actually inefficient: $(\mathbf{P}')^k$ is much less sparse than $\mathbf{P}'$

PCTL bounded until – Example


- Model check: $P_{>0.98} [F^{\leq 2} \text{ succ}] \equiv P_{>0.98} [\text{true } U^{\leq 2} \text{ succ}]$
 - $\text{Sat}(\text{true}) = S = \{s_0, s_1, s_2, s_3\}$, $\text{Sat}(\text{succ}) = \{s_3\}$
 - $S^{\text{yes}} = \{s_3\}$, $S^{\text{no}} = \emptyset$, $S^? = \{s_0, s_1, s_2\}$, $P' = P$
 - $\text{Prob}(\text{true } U^{\leq 0} \text{ succ}) = \underline{\text{succ}} = [0, 0, 0, 1]$

$$\text{Prob}(\text{true } U^{\leq 1} \text{ succ}) = P' \cdot \text{Prob}(\text{true } U^{\leq 0} \text{ succ}) = \begin{bmatrix} 0 & 1 & 0 & 0 \\ 0 & 0.01 & 0.01 & 0.98 \\ 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix} = \begin{bmatrix} 0 \\ 0.98 \\ 0 \\ 1 \end{bmatrix}$$

$$\text{Prob}(\text{true } U^{\leq 2} \text{ succ}) = P' \cdot \text{Prob}(\text{true } U^{\leq 1} \text{ succ}) = \begin{bmatrix} 0 & 1 & 0 & 0 \\ 0 & 0.01 & 0.01 & 0.98 \\ 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} 0 \\ 0.98 \\ 0 \\ 1 \end{bmatrix} = \begin{bmatrix} 0.98 \\ 0.9898 \\ 0 \\ 1 \end{bmatrix}$$

- $\text{Sat}(P_{>0.98} [F^{\leq 2} \text{ succ}]) = \{s_1, s_3\}$

The construction of P' for bounded Until




0	1	0	0
0	0.01	0.01	0.98
1	0	0	0
0	0	0	1

The probability matrix P of the DTMC on the left.

- Consider as an example to check whether the DTMC satisfies $P_{>0.99}[\text{try} \vee \neg\text{fail} \ U^{\leq 2} \text{succ}]$.
 - Calculate first the probability vector **Prob** $[\text{try} \vee \neg\text{fail} \ U^{\leq 2} \text{succ}]$.
 - From there you can calculate the set **Sat** $(P_{>0.99}[\text{try} \vee \neg\text{fail} \ U^{\leq 2} \text{succ}])$.
 - If the initial state s_0 is in the blue Sat-set** then the property $P_{>0.99}[\text{try} \vee \neg\text{fail} \ U^{\leq 2} \text{succ}]$ holds on the DTMC.

The construction of P' for bounded Until




0	1	0	0
0	0.01	0.01	0.98
1	0	0	0
0	0	0	1

The probability matrix P of the DTMC on the left.

- To calculate the probability vector $\text{Prob}[\text{try} \vee \neg\text{fail} \ U^{\leq 2} \text{succ}]$, we would like to use the matrix P above, however it will also "contain" transitions that cause you to break the green-property. So the idea is to use a "modified" matrix P' .
- We pre-calculate first the $S^{\text{yes}} = \text{Sat}(\text{succ}) = \{s3\}$. On all states in S^{yes} , you have the green property immediately (in 0 step).
- We pre-calculate S^{no} , we take $S^{\text{no}} = \text{Sat}(\neg(\text{try} \vee \neg\text{fail}) \wedge \neg \text{succ}) = \{s2\}$. Executions starting from S^{no} won't satisfy your green-property above,

The construction of P' for bounded Until



$P :$


0	1	0	0
0	0.01	0.01	0.98
1	0	0	0
0	0	0	1

$P' :$

0	1	0	0
0	0.01	0.01	0.98
0	0	0	0
0	0	0	1

1. We remove outgoing arrows from the states in S^{no} and S^{yes} .
2. We keep all arrows that go out from states which are **not** in S^{no} nor S^{yes} .
3. We add a self-loop $s \rightarrow s$ with probability 1 for any state s in S^{yes} .

Using P' for bounded Until



P' :

0	1	0	0
0	0.01	0.01	0.98
0	0	0	0
0	0	0	1

- We now use P' to iteratively calculate $\text{Prob}[\text{try } \vee \neg\text{fail } \mathbf{U}^{\leq 2} \text{succ}]$

- From S^{yes} you know that $\text{Prob}[\text{try } \vee \neg\text{fail } \mathbf{U}^{\leq 0} \text{succ}] =$

0
0
0
1

- $\text{Prob}[\text{try } \vee \neg\text{fail } \mathbf{U}^{\leq 1} \text{succ}] = P' \times$

0
0
0
1


0
0.98
0
1

- $\text{Prob}[\text{try } \vee \neg\text{fail } \mathbf{U}^{\leq 2} \text{succ}] = P' \times$

0
0.98
0
1

0.98
0.9898
0
1

So, does the property hold?



- We have calculated $\text{Prob}[\text{try} \vee \neg\text{fail} \text{ } \mathbf{U}^{\leq 2} \text{ succ}] = [0.98, 0.9898, 0, 1]$
- So, the set $\text{Sat}(P_{>0.99}[\text{try} \vee \neg\text{fail} \text{ } \mathbf{U}^{\leq 2} \text{ succ}]) = \{s_3\}$
- So we conclude that the DTMC does **not** satisfy the claimed property $P_{>0.99}[\text{try} \vee \neg\text{fail} \text{ } \mathbf{U}^{\leq 2} \text{ succ}]$.

PCTL model checking – Summary

- Computation of set $\text{Sat}(\Phi)$ for DTMC D and PCTL formula Φ
 - recursive descent of parse tree
 - combination of graph algorithms, numerical computation
- Probabilistic operator P :
 - $X \Phi$: one matrix–vector multiplication, $O(|S|^2)$
 - $\Phi_1 \cup^{\leq k} \Phi_2$: k matrix–vector multiplications, $O(k|S|^2)$
 - $\Phi_1 \cup \Phi_2$: linear equation system, at most $|S|$ variables, $O(|S|^3)$
- Complexity:
 - linear in $|\Phi|$ and polynomial in $|S|$

Overview (Part 1)

- Discrete-time Markov chains (DTMCs)
- PCTL: A temporal logic for DTMCs
- PCTL model checking
- LTL model checking
- Costs and rewards

Limitations of PCTL

- PCTL, although useful in practice, has limited expressivity
 - essentially: probability of reaching states in X , passing only through states in Y (and within k time-steps)
- More expressive logics can be used, for example:
 - LTL [Pnu77] – (non-probabilistic) linear-time temporal logic
 - PCTL* [ASB+95,BdA95] – which subsumes both PCTL and LTL
 - both allow path operators to be combined
 - (in PCTL, $P_{\sim p} [\dots]$ always contains a single temporal operator)
- Another direction: extend DTMCs with costs and rewards...

LTL – Linear temporal logic

- LTL syntax (path formulae only)

- $\psi ::= \text{true} \mid a \mid \psi \wedge \psi \mid \neg\psi \mid X\psi \mid \psi \cup \psi$
- where $a \in AP$ is an atomic proposition
- usual equivalences hold: $F\phi \equiv \text{true} \cup \phi$, $G\phi \equiv \neg(F\neg\phi)$

- LTL semantics (for a path ω)

- $\omega \models \text{true}$ always
- $\omega \models a \iff a \in L(\omega(0))$
- $\omega \models \psi_1 \wedge \psi_2 \iff \omega \models \psi_1 \text{ and } \omega \models \psi_2$
- $\omega \models \neg\psi \iff \omega \not\models \psi$
- $\omega \models X\psi \iff \omega[1\dots] \models \psi$
- $\omega \models \psi_1 \cup \psi_2 \iff \exists k \geq 0 \text{ s.t. } \omega[k\dots] \models \psi_2 \wedge \forall i < k \omega[i\dots] \models \psi_1$

where $\omega(i)$ is i^{th} state of ω , and $\omega[i\dots]$ is suffix starting at $\omega(i)$

LTL examples

- $(F \text{ tmp_fail}_1) \wedge (F \text{ tmp_fail}_2)$
 - “both servers suffer temporary failures at some point”
- $GF \text{ ready}$
 - “the server always eventually returns to a ready-state”
- $FG \text{ error}$
 - “an irrecoverable error occurs”
- $G (\text{req} \rightarrow X \text{ ack})$
 - “requests are always immediately acknowledged”

LTL for DTMCs


- Same idea as PCTL: probabilities of sets of path formulae
 - for a state s of a DTMC and an LTL formula ψ :
 - $\text{Prob}(s, \psi) = \Pr_s \{ \omega \in \text{Path}(s) \mid \omega \models \psi \}$
 - all such path sets are measurable [Var85]
- A (probabilistic) LTL specification often comprises an LTL (path) formula and a probability bound
 - e.g. $P_{\geq 1} [GF \text{ ready}]$ – “with probability 1, the server always eventually returns to a ready-state”
 - e.g. $P_{<0.01} [FG \text{ error}]$ – “with probability at most 0.01, an irrecoverable error occurs”
- PCTL* subsumes both LTL and PCTL
 - e.g. $P_{>0.5} [GF \text{ crit}_1] \wedge P_{>0.5} [GF \text{ crit}_2]$

Fundamental property of DTMCs

- Strongly connected component (SCC)
 - maximally strongly connected set of states
- Bottom strongly connected component (BSCC)
 - SCC T from which no state outside T is reachable from T

- Fundamental property of DTMCs:

- “with probability 1, a BSCC will be reached and all of its states visited infinitely often”



- Formally:

- $\Pr_s \{ \omega \in \text{Path}(s) \mid \exists i \geq 0, \exists \text{ BSCC } T \text{ such that}$
 $\forall j \geq i \ \omega(j) \in T \text{ and}$
 $\forall s' \in T \ \omega(k) = s' \text{ for infinitely many } k \} = 1$


LTL model checking for DTMCs

- LTL model checking for DTMCs relies on:
 - computing probability of reaching a set of “accepting” BSCCs
 - e.g. for two simple LTL formulae: **GF a** (“always eventually a”), **FG a** (“eventually always a”) we have:

- $\text{Prob}(s, \text{GF } a) = \text{Prob}(s, F T_{\text{GF}a})$
 - where $T_{\text{GF}a}$ = union of all BSCCs containing some state satisfying a

- $\text{Prob}(s, \text{FG } a) = \text{Prob}(s, F T_{\text{FG}a})$
 - where $T_{\text{FG}a}$ = union of all BSCCs containing only a-states

- To extend this idea to arbitrary LTL formula, we use ω -automata...



Example:

$$\begin{aligned} \text{Prob}(s_0, \text{GF } a) &= \text{Prob}(s_0, F T_{\text{GF}a}) \\ &= \text{Prob}(s_0, F \{s_3, s_2, s_5\}) \\ &= 2/3 + 1/6 = 5/6 \end{aligned}$$


Deterministic Rabin automata

- ω -automata represent sets of **infinite** words
 - e.g. Buchi automata, Rabin automata, ...
 - for probabilistic model checking, need **deterministic** automata
 - so we use deterministic Rabin automata (DRAs)
- A deterministic Rabin automaton is a tuple $(Q, \Sigma, \delta, q_0, \text{Acc})$:
 - Q is a finite set of states, $q_0 \in Q$ is an initial state
 - Σ is an alphabet, $\delta : Q \times \Sigma \rightarrow Q$ is a transition function
 - $\text{Acc} = \{ (L_i, K_i) \}_{i=1..k} \subseteq 2^Q \times 2^Q$ is an acceptance condition
- A run of a word on a DRA is accepting iff:
 - for some pair (L_i, K_i) , the states in L_i are visited finitely often and (some of) the states in K_i are visited infinitely often
 - or in LTL:
$$\bigvee_{1 \leq i \leq k} (\text{FG } \neg L_i \wedge \text{GF } K_i)$$

LTL & DRAs

- Example: DRA for **FG a**


- acceptance condition is $\text{Acc} = \{ (\{q_0\}, \{q_1\}) \}$



- Can convert any LTL formula ψ on atomic propositions AP
 - into an equivalent DRA A_ψ over alphabet 2^{AP}
 - i.e. $\omega \models \psi \Leftrightarrow \text{trace}(\omega) \in L(A_\psi)$ for any path ω
 - can potentially incur a double exponential blow-up (but, in practice, this does not occur and ψ is small anyway)
- LTL model checking for DTMCs – the basic idea
 - construct product of DTMC D and DRA A_ψ
 - compute $\text{Prob}^D(s, \psi)$ on product DTMC $D \otimes A$


Buchi vs Rabin

- Consider the LTL property $\Diamond \Box a$. This can be described by this Buchi automaton, with $\{1\}$ as the accepting state:



Notice that this Buchi is **non-deterministic**. As such, we can use it for model checking on a probabilistic model such as DTMC.

- We can however represent the property with a **deterministic** Rabin automaton, with the pair $(\{0\}, \{1\})$ as its accepting condition.



Product DTMC for a DRA

- The product DTMC $D \otimes A$ for:
 - for DTMC $D = (S, s_{init}, P, L)$ and
 - and (total) DRA $A = (Q, \Sigma, \delta, q_0, \{ (L_i, K_i) \}_{i=1..k})$
 - is the DTMC $(S \times Q, (s_{init}, q_{init}), P', L')$ where:

$$q_{init} = \delta(q_0, L(s_{init}))$$

$$P'((s_1, q_1), (s_2, q_2)) = \begin{cases} P(s_1, s_2) & \text{if } q_2 = \delta(q_1, L(s_2)) \\ 0 & \text{otherwise} \end{cases}$$

$$l_i \in L'(s, q) \text{ if } q \in L_i \text{ and } k_i \in L'(s, q) \text{ if } q \in K_i$$
- Note:
 - $D \otimes A$ can be seen as unfolding of D where q for each state (s, q) records state of automata A for path fragment so far
 - since A is deterministic, $D \otimes A$ is a DTMC
 - each path in D has a corresponding (unique) path in $D \otimes A$
 - the probabilities of paths in D are preserved in $D \otimes A$

Product DTMC for a DRA

- For DTMC D and DRA A

$$\text{Prob}^D(s, A) = \text{Prob}^{D \otimes A}((s, q_s), \bigvee_{1 \leq i \leq k} (\text{FG } \neg l_i \wedge \text{GF } k_i))$$

– where $q_s = \delta(q_0, L(s))$

- Hence:


$$\text{Prob}^D(s, A) = \text{Prob}^{D \otimes A}((s, q_s), F T_{\text{Acc}})$$

- where T_{Acc} is the union of all **accepting BSCCs** in $D \otimes A$
- an **accepting BSCC** T of $D \otimes A$ is such that, for some $1 \leq i \leq k$,
no states in T satisfy l_i and some state in T satisfies k_i


- Reduces to computing BSCCs and reachability probabilities
 - so overall complexity for LTL is doubly exponential in $|\psi|$, polynomial in $|M|$; but can be reduced to singly exponential

Example

Consider the DTMC M below. We want to check if the property $P_{>0.99}[\varphi]$ holds on M , where LTL $\varphi = \mathbf{X}(\text{try } \mathbf{W} \text{ succ})$




First, we construct a deterministic Rabin automaton that equivalently describes φ . I give a deterministic Buchi:



As a Buchi, it is also a Rabin, with accepting condition: $\{ (L, K) \}$ with $L = \emptyset$ and $K = \{v, w\}$


Intersection DTMC and Rabin

Before we construct the intersection $M \cap R$ (or product $M \otimes R$, in Kwiatkowska's notation), I prefer to first transform M so that we have labels on the arrows, by moving them to incoming arrows. The probabilities of the transitions remain the same, of course.



Intersection DTMC and Rabin


Now we can construct the intersection. This is done as we did with standard LTL, with as addition that probabilities are kept as in M:



This results in a new Rabin $R2$. Any state (s, t) in $R2$ with t in R 's L is marked as L in $R2$. If t was in R 's K , we mark (s, t) as K in $R2$. So, for above example we get L is empty, and all the blue states are K .

Calculating $\text{Prob}(\varphi)$

So we have this $R2 = M \cap R$, with $L = \emptyset$ and $K = \{ \text{the blue states} \}$. Any execution satisfying φ can only pass L finitely many often and should pass one of K infinitely many often.



The colored boxes are the SCCs of $R2$. Only the blue one, let's call it Z , is a bottom SCC. Z contains a member of K , and no member of Z is in L . So, any execution satisfying φ will end up in Z .

We know:

- $\text{Prob}((s_3, w), \varphi) = 1$
- $\text{Prob}((z, u), \varphi) = \text{Prob}((s_0, v), \varphi) = \text{Prob}((s_1, v), \varphi)$
- $\text{Prob}((s_1, v), \varphi) = 0.01 * \text{Prob}((s_1, v), \varphi) + 0.98 * \text{Prob}((s_3, w), \varphi)$


Solving the equations gives $\text{Prob}((z, u), \varphi) = \text{Prob}((s_1, v), \varphi) = 98/99 < 0.99$.

So, the answer to the original question if $P_{>0.99}[\varphi]$ holds is a "no".


Another Example: LTL for DTMCs

- Compute $\text{Prob}(s_0, G\neg b \wedge GF a)$ for DTMC **D**:

DTMC **D**




DRA A_ψ for $\psi = G\neg b \wedge GF a$




Another Example: LTL for DTMCs


DTMC **D**



DRA A_ψ for $\psi = \mathbf{G}\neg b \wedge \mathbf{GF} a$




Product DTMC $\mathbf{D} \otimes A_\psi$




Another Example: LTL for DTMCs


DTMC **D**



DRA A_ψ for $\psi = G\neg b \wedge GF a$



Product DTMC $D \otimes A_\psi$



$$\begin{aligned} \text{Prob}^D(s, \psi) &= \text{Prob}^{D \otimes A_\psi}(F T_1) \\ &= 3/4. \end{aligned}$$

Overview (Part 1)

- Discrete-time Markov chains (DTMCs)
- PCTL: A temporal logic for DTMCs
- PCTL model checking
- LTL model checking
- Costs and rewards

Costs and rewards

- We augment DTMCs with rewards (or, conversely, costs)
 - real-valued quantities assigned to states and/or transitions
 - these can have a wide range of possible interpretations
- Some examples:
 - elapsed time, power consumption, size of message queue, number of messages successfully delivered, net profit, ...
- Costs? or rewards?
 - mathematically, no distinction between rewards and costs
 - when interpreted, we assume that it is desirable to minimise costs and to maximise rewards
 - we will consistently use the terminology “rewards” regardless

Reward-based properties


- Properties of DTMCs augmented with rewards
 - allow a wide range of quantitative measures of the system
 - basic notion: expected value of rewards
 - formal property specifications will be in an extension of PCTL
- More precisely, we use two distinct classes of property...
- **Instantaneous** properties
 - the expected value of the reward at some time point
- **Cumulative** properties
 - the expected cumulated reward over some period

DTMC reward structures

- For a DTMC $(S, s_{\text{init}}, \mathbf{P}, L)$, a reward structure is a pair $(\underline{r}, \underline{t})$
 - $\underline{r} : S \rightarrow \mathbb{R}_{\geq 0}$ is the **state reward function** (vector)
 - $\underline{t} : S \times S \rightarrow \mathbb{R}_{\geq 0}$ is the **transition reward function** (matrix)
- Example (for use with instantaneous properties)
 - “size of message queue”: $\underline{r}$ maps each state to the number of jobs in the queue in that state, $\underline{t}$ is not used
- Examples (for use with cumulative properties)
 - “**time-steps**”: $\underline{r}$ returns 1 for all states and $\underline{t}$ is zero (equivalently, $\underline{r}$ is zero and $\underline{t}$ returns 1 for all transitions)
 - “**number of messages lost**”: $\underline{r}$ is zero and $\underline{t}$ maps transitions corresponding to a message loss to 1
 - “**power consumption**”: $\underline{r}$ is defined as the per-time-step energy consumption in each state and $\underline{t}$ as the energy cost of each transition

PCTL and rewards

- Extend PCTL to incorporate reward-based properties
 - add an R operator, which is similar to the existing P operator



– where $r \in \mathbb{R}_{\geq 0}$, $\sim \in \{<, >, \leq, \geq\}$, $k \in \mathbb{N}$

- $R_{\sim r}[\cdot]$ means “the **expected value** of $\cdot$ satisfies $\sim r$ ”

Types of reward formulas

- **Instantaneous:** $R_{\sim r} [I^k]$
 - “the expected value of the state reward at time-step k is $\sim r$ ”
 - e.g. “the expected queue size after exactly 90 seconds”
- **Cumulative:** $R_{\sim r} [C^{\leq k}]$
 - “the expected reward cumulated up to time-step k is $\sim r$ ”
 - e.g. “the expected power consumption over one hour”
- **Reachability:** $R_{\sim r} [F \phi]$
 - “the expected reward cumulated before reaching a state satisfying ϕ is $\sim r$ ”
 - e.g. “the expected time for the algorithm to terminate”

Reward formula semantics

- Formal semantics of the three reward operators
 - based on random variables over (infinite) paths
- Recall:
 - $s \models P_{\sim p} [\psi] \Leftrightarrow \Pr_s \{ \omega \in \text{Path}(s) \mid \omega \models \psi \} \sim p$
- For a state s in the DTMC:
 - $s \models R_{\sim r} [I^k] \Leftrightarrow \text{Exp}(s, X_{I^k}) \sim r$
 - $s \models R_{\sim r} [C^{\leq k}] \Leftrightarrow \text{Exp}(s, X_{C^{\leq k}}) \sim r$
 - $s \models R_{\sim r} [F \Phi] \Leftrightarrow \text{Exp}(s, X_{F\Phi}) \sim r$

where: $\text{Exp}(s, X)$ denotes the **expectation** of the **random variable** $X : \text{Path}(s) \rightarrow \mathbb{R}_{\geq 0}$ with respect to the **probability measure** $\Pr_s$

Reward formula semantics

- Definition of random variables:
 - for an infinite path $\omega = s_0 s_1 s_2 \dots$

$$X_{I=k}(\omega) = \underline{\rho}(s_k)$$

$$X_{C \leq k}(\omega) = \begin{cases} 0 & \text{if } k = 0 \\ \sum_{i=0}^{k-1} \underline{\rho}(s_i) + \iota(s_i, s_{i+1}) & \text{otherwise} \end{cases}$$

$$X_{F\phi}(\omega) = \begin{cases} 0 & \text{if } s_0 \in \text{Sat}(\phi) \\ \infty & \text{if } s_i \notin \text{Sat}(\phi) \text{ for all } i \geq 0 \\ \sum_{i=0}^{k_\phi-1} \underline{\rho}(s_i) + \iota(s_i, s_{i+1}) & \text{otherwise} \end{cases}$$

- where $k_\phi = \min\{j \mid s_j \models \phi\}$

Model checking reward properties

- Instantaneous: $R_{\sim r} [I^k]$
- Cumulative: $R_{\sim r} [C^{\leq t}]$
 - variant of the method for computing bounded until probabilities
 - solution of **recursive equations**
- Reachability: $R_{\sim r} [F \phi]$
 - similar to computing until probabilities
 - precomputation phase (identify infinite reward states)
 - then reduces to solving a **system of linear equation**
- For more details, see e.g. [\[KNP07a\]](#)

Summary

- Probabilistic model checking
 - automated quantitative verification of stochastic systems
 - to model randomisation, failures, ...
- Discrete-time Markov chains (DTMCs)
 - state transition systems + discrete probabilistic choice
 - probability space over paths through a DTMC
- Property specifications
 - probabilistic extensions of temporal logic, e.g. PCTL, LTL
 - also: expected value of costs/rewards
- Model checking algorithms
 - combination of graph-based algorithms, numerical computation, automata constructions
- Tomorrow: Markov decision processes (MDPs)