



Probabilistic Model Checking

Marta Kwiatkowska
Dave Parker

Oxford University Computing Laboratory

ESSLLI'10 Summer School, Copenhagen, August 2010



Part 2

Markov decision processes

Overview (Part 2)


- Markov decision processes (MDPs)
- Adversaries & probability spaces
- PCTL for MDPs
- PCTL model checking
- Further model checking (LTL, costs & rewards)
- Case study: Firewire root contention

Nondeterminism

- Some aspects of a system may not be probabilistic and should not be modelled probabilistically; for example:
- **Concurrency** – scheduling of parallel components
 - e.g. randomised distributed algorithms – multiple probabilistic processes operating **asynchronously**
- **Underspecification** – unknown model parameters
 - e.g. a probabilistic communication protocol designed for message propagation delays of between d_{\min} and d_{\max}
- **Unknown environments**
 - e.g. probabilistic security protocols – unknown adversary

Markov decision processes

- Markov decision processes (MDPs)
 - extension of DTMCs which allow **nondeterministic choice**
- Like DTMCs:
 - discrete set of states representing possible configurations of the system being modelled
 - transitions between states occur in discrete time-steps
- Probabilities and nondeterminism
 - in each state, a nondeterministic choice between several discrete probability distributions over successor states



Markov decision processes

- Formally, an MDP M is a tuple $(S, s_{\text{init}}, \text{Steps}, L)$ where:
 - S is a finite set of states (“state space”)
 - $s_{\text{init}} \in S$ is the initial state
 - Steps** : $S \rightarrow 2^{\text{Act} \times \text{Dist}(S)}$ is the **transition probability function** where Act is a set of actions and $\text{Dist}(S)$ is the set of discrete probability distributions over the set S
 - $L : S \rightarrow 2^{\text{AP}}$ is a labelling with atomic propositions

For example:

$\text{Steps}(s_0) = \{ (a, u) \}$ where u is a “distribution” describing the probability of getting to a state if we take an a -transition on s_1 :

$$u(s_0)=0, u(s_1)=1, u(s_2)=0, u(s_3)=0$$


For state s_1 , we have

$$\text{Steps}(s_1) = \{ (b, v), (c, w) \}$$

where v and w are distributions :

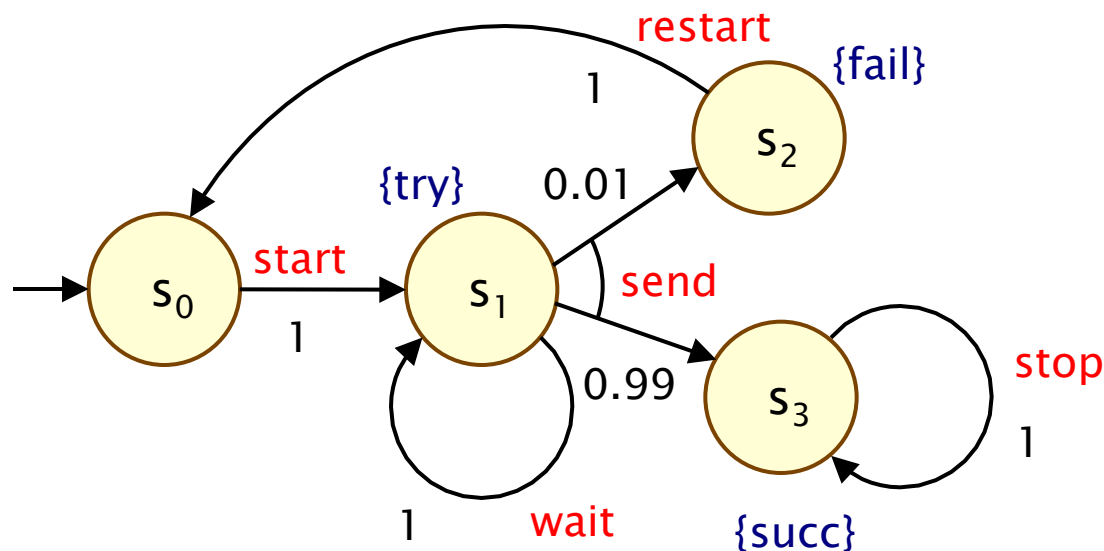
$$v(s_0) = 0.7, v(s_1) = 0.3, v(s_2) = 0, v(s_3) = 0$$

$$w(s_0) = 0, w(s_1) = 0, w(s_2) = 0.5, w(s_3) = 0.5$$



Simple MDP example



- Modification of the simple DTMC communication protocol
 - after one step, process **starts** trying to send a message
 - then, a nondeterministic choice between: (a) **waiting** a step because the channel is unready; (b) **sending** the message
 - if the latter, with probability 0.99 send **successfully** and **stop**
 - and with probability 0.01, message sending **fails**, **restart**



Example – Parallel composition

Asynchronous parallel composition of two 3-state DTMCs

Action labels omitted here



at every state above there is a non-deterministic choice between taking a blue or a brown transition.

Paths and probabilities

- A (finite or infinite) path through an MDP
 - is a sequence of states and action/distribution pairs
 - e.g. $s_0(a_0, \mu_0)s_1(a_1, \mu_1)s_2\dots$
 - such that $(a_i, \mu_i) \in \text{Steps}(s_i)$ and $\mu_i(s_{i+1}) > 0$ for all $i \geq 0$
 - represents an **execution** (i.e. one possible behaviour) of the system which the MDP is modelling
 - note that a **path resolves both types of choices**: nondeterministic and probabilistic
- To consider the probability of some behaviour of the MDP
 - first need to **resolve the nondeterministic choices**
 - ...which results in a **DTMC**
 - ...for which we can define a **probability measure over paths**

Adversaries


- An **adversary** resolves nondeterministic choice in an MDP
 - also known as “schedulers”, “strategies” or “policies”
- Formally:
 - an adversary A of an MDP M is a function **mapping** every **finite path** $\omega = s_0(a_1, \mu_1)s_1 \dots s_n$ to an **element of $\text{Steps}(s_n)$**
- For each A can define a probability measure Pr_s^A over paths
 - constructed through an **infinite state DTMC** $(\text{Path}_{\text{fin}}^A(s), s, \mathbf{P}_s^A)$
 - **states** of the DTMC are the **finite paths of A starting in state s**
 - initial state is s (the path starting in s of length 0)
 - $\mathbf{P}_s^A(\omega, \omega') = \mu(s)$ if $\omega' = \omega(a, \mu)s$ and $A(\omega) = (a, \mu)$
 - $\mathbf{P}_s^A(\omega, \omega') = 0$ otherwise

Adversaries – Examples

- Consider the simple MDP below
 - note that s_1 is the only state for which $|\text{Steps}(s)| > 1$
 - i.e. s_1 is the only state for which an adversary makes a choice
 - let μ_b and μ_c denote the probability distributions associated with actions b and c in state s_1

- Adversary A_1

- picks action c the first time
- $A_1(s_0s_1) = (c, \mu_c)$





- Adversary A_2

- picks action b the first time, then c
- $A_2(s_0s_1) = (b, \mu_b)$, $A_2(s_0s_1s_1) = (c, \mu_c)$, $A_2(s_0s_1s_0s_1) = (c, \mu_c)$

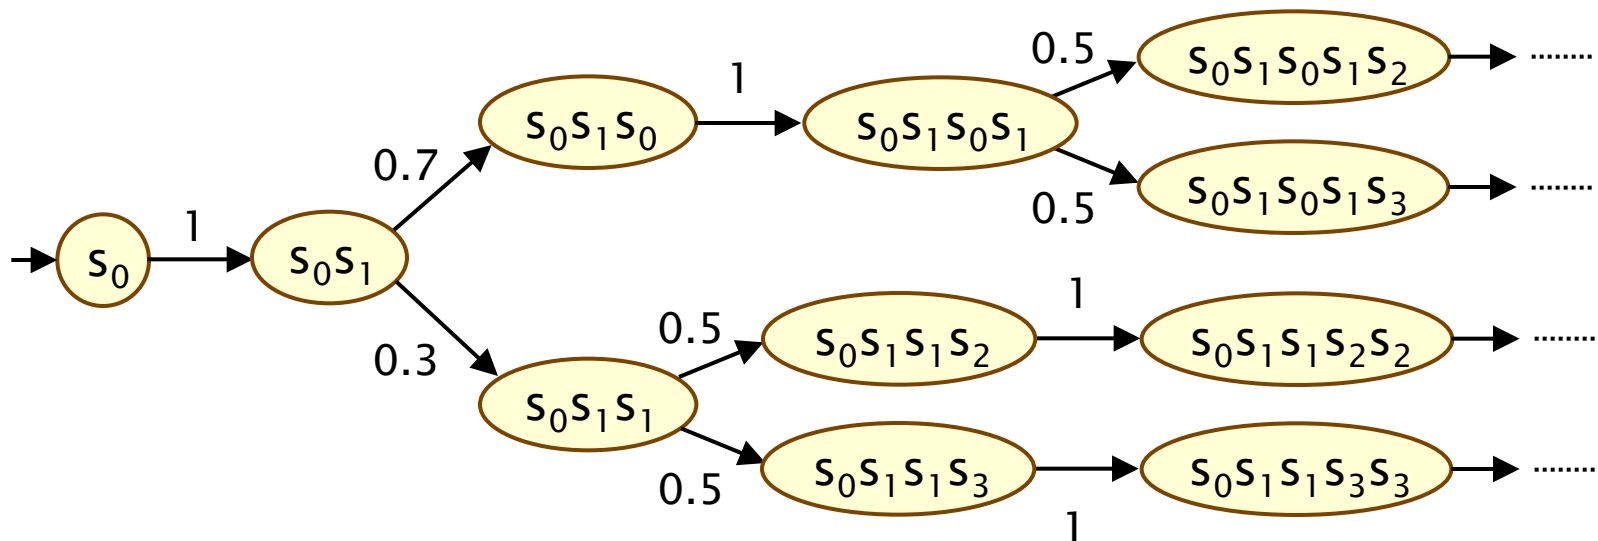

Adversaries – Examples

- Fragment of DTMC for adversary A_1
 - A_1 picks action c the first time




Adversaries – Examples

- Fragment of DTMC for adversary A_2
 - A_2 picks action b, then c



Memoryless adversaries

- **Memoryless adversaries** always pick same choice in a state
 - also known as: positional, Markov, simple
 - formally, for adversary A :
 - $A(s_0(a_1, \mu_1) s_1 \dots s_n)$ depends only on s_n
 - resulting DTMC can be mapped to a $|S|$ -state DTMC
- From previous example:
 - adversary A_1 (picks c in s_1) is memoryless, A_2 is not



Overview (Part 2)

- Markov decision processes (MDPs)
- Adversaries & probability spaces
- PCTL for MDPs
- PCTL model checking
- Further model checking (LTL, costs & rewards)
- Case study: Firewire root contention

PCTL for MDPs

- The temporal logic PCTL can also describe MDP properties
- Identical syntax to the DTMC case:

– $\phi ::= \text{true} \mid a \mid \phi \wedge \phi \mid \neg \phi \mid P_{\sim p} [\psi]$ (state formulas)

– $\psi ::= X \phi \mid \phi U^{\leq k} \phi \mid \phi U \phi$ (path formulas)

“next”

“bounded
until”


“until”

ψ is true with
probability $\sim p$

- Semantics are also the same as DTMCs for:
 - atomic propositions, logical operators, path formulas

PCTL semantics for MDPs

- Semantics of the probabilistic operator P
 - can only define **probabilities** for a **specific adversary A**
 - $s \models P_{\sim p} [\psi]$ means “the probability, from state s , that ψ is true for an outgoing path satisfies $\sim p$ **for all adversaries A** ”
 - formally $s \models P_{\sim p} [\psi] \Leftrightarrow \text{Prob}^A(s, \psi) \sim p$ for all adversaries A
 - where $\text{Prob}^A(s, \psi) = \Pr^A_s \{ \omega \in \text{Path}^A(s) \mid \omega \models \psi \}$



$$\text{Prob}^A(s, \psi) \sim p$$

Minimum and maximum probabilities

- Letting:
 - $p_{\max}(s, \psi) = \sup_A \text{Prob}^A(s, \psi)$
 - $p_{\min}(s, \psi) = \inf_A \text{Prob}^A(s, \psi)$
- We have:
 - if $\sim \in \{\geq, >\}$, then $s \models P_{\sim p} [\psi] \iff p_{\min}(s, \psi) \sim p$
 - if $\sim \in \{<, \leq\}$, then $s \models P_{\sim p} [\psi] \iff p_{\max}(s, \psi) \sim p$
- Model checking $P_{\sim p} [\psi]$ reduces to the computation over all adversaries of either:
 - the **minimum probability** of ψ holding
 - the **maximum probability** of ψ holding
- Crucial result for model checking PCTL on MDPs
 - memoryless adversaries suffice, i.e. there are always memoryless adversaries A_{\min} and A_{\max} for which:
 - $\text{Prob}^{A_{\min}}(s, \psi) = p_{\min}(s, \psi)$ and $\text{Prob}^{A_{\max}}(s, \psi) = p_{\max}(s, \psi)$

Quantitative properties

- For PCTL properties with P as the outermost operator
 - quantitative form (two types): $P_{\min=?} [\psi]$ and $P_{\max=?} [\psi]$
 - i.e. “**what is the minimum/maximum probability (over all adversaries) that path formula ψ is true?**”
 - corresponds to an analysis of **best-case** or **worst-case** behaviour of the system
 - model checking is no harder since compute the values of $p_{\min}(s, \psi)$ or $p_{\max}(s, \psi)$ anyway
 - useful to spot patterns/trends
- Example: CSMA/CD protocol
 - “min/max probability that a message is sent within the deadline”

Some real PCTL examples

- Byzantine agreement protocol
 - $P_{\min=?} [F (\text{agreement} \wedge \text{rounds} \leq 2)]$
 - “what is the minimum probability that agreement is reached within two rounds?”
- CSMA/CD communication protocol
 - $P_{\max=?} [F \text{ collisions} = k]$
 - “what is the maximum probability of k collisions?”
- Self-stabilisation protocols
 - $P_{\min=?} [F^{\leq t} \text{ stable}]$
 - “what is the minimum probability of reaching a stable state within k steps?”

Overview (Part 2)

- Markov decision processes (MDPs)
- Adversaries & probability spaces
- PCTL for MDPs
- **PCTL model checking**
- Further model checking (LTL, costs & rewards)
- Case study: Firewire root contention

PCTL model checking for MDPs

- Algorithm for PCTL model checking [BdA95]
 - inputs: MDP $M=(S, s_{init}, \text{Steps}, L)$, PCTL formula ϕ
 - output: $\text{Sat}(\phi) = \{ s \in S \mid s \models \phi \}$ = set of states satisfying ϕ
- Basic algorithm same as PCTL model checking for DTMCs
 - proceeds by induction on parse tree of ϕ
 - non-probabilistic operators (true , a , \neg , \wedge) straightforward
- Only need to consider $P_{\sim p} [\psi]$ formulas
 - reduces to computation of $p_{\min}(s, \psi)$ or $p_{\max}(s, \psi)$ for all $s \in S$
 - dependent on whether $\sim \in \{\geq, >\}$ or $\sim \in \{<, \leq\}$
 - these slides cover the case $p_{\min}(s, \phi_1 \cup \phi_2)$, i.e. $\sim \in \{\geq, >\}$
 - case for maximum probabilities is very similar
 - next ($X \phi$) and bounded until ($\phi_1 \cup^{\leq k} \phi_2$) are straightforward extensions of the DTMC case


PCTL next for MDPs

- Computation of probabilities for PCTL next operator
- Consider case of minimum probabilities...

- $\text{Sat}(P_{\sim p}[X \phi]) = \{ s \in S \mid p_{\min}(s, X \phi) \sim p \}$
- need to compute $p_{\min}(s, X \phi)$ for all $s \in S$

- Recall in the DTMC case

- sum outgoing probabilities for transitions to ϕ -states
- $\text{Prob}(s, X \phi) = \sum_{s' \in \text{Sat}(\phi)} P(s, s')$



- For MDPs, perform computation for **each distribution** available in s and then **take minimum**:


- $p_{\min}(s, X \phi) = \min \{ \sum_{s' \in \text{Sat}(\phi)} \mu(s') \mid (a, \mu) \in \text{Steps}(s) \}$

- Maximum probabilities case is analogous

PCTL next – Example

- Model check: $P_{\geq 0.5} [X \text{ heads}]$
 - lower probability bound so **minimum probabilities** required
 - $\text{Sat}(\text{heads}) = \{s_2\}$
 - e.g. $p_{\min}(s_1, X \text{ heads}) = \min(0, 0.5) = 0$
 - can do all at once with matrix–vector multiplication:

$$\text{Steps} \cdot \underline{\text{heads}} = \begin{bmatrix} 0 & 1 & 0 & 0 \\ 0.7 & 0.3 & 0 & 0 \\ 0 & 0 & 0.5 & 0.5 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0.5 \\ 1 \\ 0 \end{bmatrix}$$




- Extracting the minimum for each state yields
 - $\underline{p}_{\min}(X \text{ heads}) = [0, 0, 1, 0]$
 - $\text{Sat}(P_{\geq 0.5} [X \text{ heads}]) = \{s_2\}$

PCTL until for MDPs

- Computation of probabilities $p_{\min}(s, \phi_1 \text{ U } \phi_2)$ for all $s \in S$
- First identify all states where the **probability** is **1** or **0**
 - “precomputation” algorithms, yielding sets $S^{\text{yes}}, S^{\text{no}}$
- Then compute (min) probabilities for remaining states ($S^?$)
 - either: solve linear programming problem
 - or: approximate with an iterative solution method
 - or: use policy iteration

Example:
 $P_{\geq p} [F a]$
 \equiv
 $P_{\geq p} [\text{true U } a]$




PCTL until – Precomputation

- Identify all states where $p_{\min}(s, \phi_1 \cup \phi_2)$ is 1 or 0
 - $S^{\text{yes}} = \text{Sat}(P_{\geq 1} [\phi_1 \cup \phi_2])$, $S^{\text{no}} = \text{Sat}(\neg P_{>0} [\phi_1 \cup \phi_2])$
- Two graph-based precomputation algorithms:
 - algorithm Prob1A computes S^{yes}
 - for all adversaries the probability of satisfying $\phi_1 \cup \phi_2$ is 1
 - algorithm Prob0E computes S^{no}
 - there exists an adversary for which the probability is 0

Example:

$P_{\geq p} [F a]$



$$S^{\text{yes}} = \text{Sat}(P_{\geq 1} [F a])$$

$$S^{\text{no}} = \text{Sat}(\neg P_{>0} [F a])$$

Calculating $\text{pmin}(s, \varphi_1 \mathbf{U} \varphi_2)$

$\text{pmin}(s, \varphi)$ = the **minimum** probability for having executions starting from s satisfying φ , regardless the adversary. To calculate $\text{pmin}(s, \varphi_1 \mathbf{U} \varphi_2)$:

1. Calculate first the sets S^{yes} and S^{no} .
 - $S^{\text{yes}} = \{ s \mid P(s, \varphi_1 \mathbf{U} \varphi_2) \geq 1, \text{ for } \mathbf{all} \text{ adversaries} \} \rightarrow$ with algorithm prob1A.
 - $S^{\text{no}} = \{ s \mid P(s, \varphi_1 \mathbf{U} \varphi_2) \leq 0, \text{ for } \mathbf{some} \text{ adversary} \} \rightarrow$ with algorithm prob0E.
2. For any state s in S^{yes} , we then know that $\text{pmin}(s, \varphi_1 \mathbf{U} \varphi_2) \geq 1$.
3. For any state s in S^{no} we have $\text{pmin}(s, \varphi_1 \mathbf{U} \varphi_2) \leq 0$.
4. We then proceed with calculating the pmin for the remaining states (which are not in S^{yes} nor S^{no}).

Algorithm Prob0E

- The algorithm below first calculates the set R of all states s satisfying $E[s, \varphi_1 \cup \varphi_2]$, regardless the adversary. So, for any state in R , and for any adversary $\text{Prob}(s, \varphi_1 \cup \varphi_2) > 0$.
- S^{no} is just complement S/R .
- $\text{Sat}(\varphi_1)$ and $\text{Sat}(\varphi_2)$ in the paremeters are the set of states on which φ_1 and φ_2 respectively hold.

$\text{PROB0E}(\text{Sat}(\phi_1), \text{Sat}(\phi_2))$

```
1.   $R := \text{Sat}(\phi_2)$ 
2.   $done := \text{false}$ 
3.  while ( $done = \text{false}$ )
4.       $R' := R \cup \{s \in \text{Sat}(\phi_1) \mid \forall \mu \in \text{Steps}(s). \exists s' \in R. \mu(s') > 0\}$ 
5.      if ( $R' = R$ ) then  $done := \text{true}$ 
6.       $R := R'$ 
7.  endwhile
8.  return  $S \setminus R$ 
```

Prob1A

- Calculate first the set F of states from where we have a path passing exclusively through states satisfying $\varphi_1 \wedge \neg\varphi_2$ and ends in S^{no} , under **some** adversary.
By definition this F also includes S^{no} .
- So any state s in F has $\text{Prob}(s, \varphi_1 \mathbf{U} \varphi_2) < 1$, for some adversary. In other words $\text{pmin}(s, \varphi_1 \mathbf{U} \varphi_2) < 1$.
- P^{yes} in the complement S/F .
- **Note:** for the calculation of $\text{pmin}(s, \varphi_1 \mathbf{U} \varphi_2)$, we can also just take $S^{\text{yes}} = \text{Sat}(\varphi_2)$. The calculation would still work, though it would take more steps to get its final results.


Method 1 – Linear programming

- Probabilities $p_{\min}(s, \phi_1 \cup \phi_2)$ for remaining states in the set $S^? = S \setminus (S^{\text{yes}} \cup S^{\text{no}})$ can be obtained as the unique solution of the following **linear programming (LP)** problem:

$$\begin{aligned} &\text{maximize } \sum_{s \in S^?} x_s \text{ subject to the constraints :} \\ &\quad x_s \leq \sum_{s' \in S^?} \mu(s') \cdot x_{s'} + \sum_{s' \in S^{\text{yes}}} \mu(s') \\ &\quad \text{for all } s \in S^? \text{ and for all } (a, \mu) \in \text{Steps}(s) \end{aligned}$$

- Simple case of a more general problem known as the **stochastic shortest path problem** [BT91]
- This can be solved with standard techniques
 - e.g. Simplex, ellipsoid method, branch-and-cut

Example – PCTL until (LP)



Let $x_i = p_{\min}(s_i, F a)$


S^{yes} : $x_2=1$, S^{no} : $x_3=0$

For $S^? = \{x_0, x_1\}$:

Maximise x_0+x_1 subject to constraints:

- $x_0 \leq x_1$
- $x_0 \leq 0.25 \cdot x_0 + 0.5$
- $x_1 \leq 0.1 \cdot x_0 + 0.5 \cdot x_1 + 0.4$

Example – PCTL until (LP)






Let $x_i = p_{\min}(s_i, F a)$

S_{yes} : $x_2=1$, S_{no} : $x_3=0$


For $S^? = \{x_0, x_1\}$:

Maximise x_0+x_1 subject to constraints:

- $x_0 \leq x_1$
- $x_0 \leq 2/3$
- $x_1 \leq 0.2 \cdot x_0 + 0.8$



Example – PCTL until (LP)




Let $x_i = p_{\min}(s_i, F a)$

S^{yes} : $x_2=1$, S^{no} : $x_3=0$

For $S^? = \{x_0, x_1\}$:

Maximise x_0+x_1 subject to constraints:

- $x_0 \leq x_1$
- $x_0 \leq 2/3$
- $x_1 \leq 0.2 \cdot x_0 + 0.8$




Solution:

(x_0, x_1)

=

$(2/3, 14/15)$

Example – PCTL until (LP)




Let $x_i = p_{\min}(s_i, F a)$

S^{yes} : $x_2=1$, S^{no} : $x_3=0$

For $S^? = \{x_0, x_1\}$:

Maximise x_0+x_1 subject to constraints:

- $x_0 \leq x_1$
- $x_0 \leq 2/3$
- $x_1 \leq 0.2 \cdot x_0 + 0.8$



Method 2 – Value iteration


- For probabilities $p_{\min}(s, \phi_1 \cup \phi_2)$ it can be shown that:

– $p_{\min}(s, \phi_1 \cup \phi_2) = \lim_{n \rightarrow \infty} x_s^{(n)}$ where:

$$x_s^{(n)} = \begin{cases} 1 & \text{if } s \in S^{\text{yes}} \\ 0 & \text{if } s \in S^{\text{no}} \\ 0 & \text{if } s \in S^? \text{ and } n = 0 \\ \min_{(a, \mu) \in \text{Steps}(s)} \left(\sum_{s' \in S} \mu(s') \cdot x_{s'}^{(n-1)} \right) & \text{if } s \in S^? \text{ and } n > 0 \end{cases}$$

- This forms the basis for an (approximate) iterative solution
 - iterations terminated when solution converges sufficiently

Example – PCTL until (value iteration)



Compute: $p_{\min}(s_i, F a)$

$S^{yes} = \{x_2\}$, $S^{no} = \{x_3\}$, $S^? = \{x_0, x_1\}$

$[x_0^{(n)}, x_1^{(n)}, x_2^{(n)}, x_3^{(n)}]$


$n=0: [0, 0, 1, 0]$

$n=1: [\min(0, 0.25 \cdot 0 + 0.5),$
 $0.1 \cdot 0 + 0.5 \cdot 0 + 0.4, 1, 0]$
 $= [0, 0.4, 1, 0]$

$n=2: [\min(0.4, 0.25 \cdot 0 + 0.5),$
 $0.1 \cdot 0 + 0.5 \cdot 0.4 + 0.4, 1, 0]$
 $= [0.4, 0.6, 1, 0]$


$n=3: \dots$

Example – PCTL until (value iteration)



	$[x_0^{(n)}, x_1^{(n)}, x_2^{(n)}, x_3^{(n)}]$
n=0:	$[0.000000, 0.000000, 1, 0]$
n=1:	$[0.000000, 0.400000, 1, 0]$
n=2:	$[0.400000, 0.600000, 1, 0]$
n=3:	$[0.600000, 0.740000, 1, 0]$
n=4:	$[0.650000, 0.830000, 1, 0]$
n=5:	$[0.662500, 0.880000, 1, 0]$
n=6:	$[0.665625, 0.906250, 1, 0]$
n=7:	$[0.666406, 0.919688, 1, 0]$
n=8:	$[0.666602, 0.926484, 1, 0]$
n=9:	$[0.666650, 0.929902, 1, 0]$
...	
n=20:	$[0.666667, 0.933332, 1, 0]$
n=21:	$[0.666667, 0.933332, 1, 0]$
	$\approx [2/3, 14/15, 1, 0]$

Example – Value iteration + LP



$[x_0^{(n)}, x_1^{(n)}, x_2^{(n)}, x_3^{(n)}]$

n=0: [0.000000, 0.000000, 1, 0]

n=1: [0.000000, 0.400000, 1, 0]

n=2: [0.400000, 0.600000, 1, 0]

n=3: [0.600000, 0.740000, 1, 0]

n=4: [0.650000, 0.830000, 1, 0]

n=5: [0.662500, 0.880000, 1, 0]

n=6: [0.665625, 0.906250, 1, 0]

n=7: [0.666406, 0.919688, 1, 0]

n=8: [0.666602, 0.926484, 1, 0]

n=9: [0.666650, 0.929902, 1, 0]

...

n=20: [0.666667, 0.933332, 1, 0]

n=21: [0.666667, 0.933332, 1, 0]

$\approx [2/3, 14/15, 1, 0]$

PCTL until for MDPs – Prob0A

- Maximum probabilities 0
 - $S^{\text{no}} = \{ s \in S \mid p_{\max}(s, \phi_1 \cup \phi_2) = 0 \}$


$\text{PROB0A}(\text{Sat}(\phi_1), \text{Sat}(\phi_2))$
<pre>1. $R := \text{Sat}(\phi_2)$ 2. $done := \text{false}$ 3. while ($done = \text{false}$) 4. $R' := R \cup \{ s \in \text{Sat}(\phi_1) \mid \exists \mu \in \text{Steps}(s) . \exists s' \in R . \mu(s') > 0 \}$ 5. if ($R' = R$) then $done := \text{true}$ 6. $R := R'$ 7. endwhile 8. return $S \setminus R$</pre>

PCTL until for MDPs – Prob1E

- Maximum probabilities 1
 - $S^{\text{yes}} = \{ s \in S \mid p_{\max}(s, \phi_1 \cup \phi_2) = 1 \} = \text{Sat}(\neg P_{<1} [\phi_1 \cup \phi_2])$
- Prob1E algorithm (see next slide)
 - two nested loops (double fixed point)
 - result, stored in R, will be S^{yes} ; initially R is S
 - iteratively remove (some) states u with $p_{\max}(u, \phi_1 \cup \phi_2) < 1$
 - i.e. remove (some) states for which, under no adversary σ , is $\text{Prob}^\sigma(s, \phi_1 \cup \phi_2) = 1$
 - done by inner loop which computes subset R' of R
 - R' contains ϕ_1 -states with a probability distribution for which all transitions stay within R and at least one eventually reaches ϕ_2
 - note: after first iteration, R contains:
 - $\{ s \mid \text{Prob}^A(s, \phi_1 \cup \phi_2) > 0 \text{ for some } A \}$
 - essentially: execution of Prob0A and removal of S^{no} from R

Prob1 E – Example

- $S^{\text{yes}} = \{ s \in S \mid p_{\max}(s, \neg a \cup b) = 1 \}$
- $R = \{ 0, 1, 2, 3, 4, 5, 6 \}$
 - $R' = \{2\}$; $R' = \{1, 2, 5\}$; $R' = \{1, 2, 4, 5\}$; $R' = \{1, 2, 4, 5, 6\}$
- $R = \{ 1, 2, 4, 5, 6 \}$
 - $R' = \{2\}$; $R' = \{1, 2, 5\}$
- $R = \{ 1, 2, 5 \}$
 - $R' = \{2\}$; $R' = \{1, 2, 5\}$
- $R = \{ 1, 2, 5 \}$
- $S^{\text{yes}} = \{ 1, 2, 5 \}$



PCTL model checking – Summary

- Computation of set $\text{Sat}(\Phi)$ for MDP M and PCTL formula Φ
 - recursive descent of parse tree
 - combination of graph algorithms, numerical computation
- Probabilistic operator P :
 - $X \Phi$: one matrix–vector multiplication, $O(|S|^2)$
 - $\Phi_1 \cup^{\leq k} \Phi_2$: k matrix–vector multiplications, $O(k|S|^2)$
 - $\Phi_1 \cup \Phi_2$: linear programming problem, polynomial in $|S|$ (assuming use of linear programming)
- Complexity:
 - linear in $|\Phi|$ and polynomial in $|S|$
 - S is states in MDP, assume $|\text{Steps}(s)|$ is constant

Summary

- Markov decision processes (MDPs)
 - extend DTMCs with nondeterminism
 - to model concurrency, underspecification, ...
- An adversary resolve nondeterminism in an MDP
 - induce a probability space over paths
 - consider minimum/maximum probabilities over all adversaries
- Property specifications
 - use e.g. PCTL or LTL, as for DTMCs
 - but quantify over all adversaries
- Model checking algorithms
 - covered three basic techniques for MDPs: linear programming, value iteration, or policy iteration
- ~~Tomorrow: continuous time Markov chains (CTMCs)~~