

Probabilistic Model Checking

Wishnu Prasetya


wishnu@cs.uu.nl

www.cs.uu.nl/docs/vakken/pv

additional slides DTMC

Probability of taking a path or a set of paths

A “DTMC” :




- Consider a path ω e.g. s_0, s_1, s_2, s_0 . The probability that the system follows this path when executed with the starting state s_0 is denoted by $P_{s_0}(\omega)$. Or simply $P(\omega)$ if it is clear which s_0 is meant. It is the product of the probability of each transition in ω .

Example: for the above ω , $P(\omega) = 1 * 0.01 * 1 = 0.01$

- For a **set of paths** U (starting from s_0), the probability that the system's execution follows **one of** the paths in U , denoted by $P(U)$, is $\sum_{\omega \in U} P(\omega)$.

Probability of taking a path or a set of paths




Example: consider U = the set of paths that ends in s_2 . Note that U is infinite: $U = \{ 02, 0102, 010102, \dots \}$. But we can calculate $P(U)$.

$$P(U) = 0.5 + 0.5^2 + 0.5^3 + \dots = \sum_{k \geq 0} 0.5^k$$

$$= 0.5 * \frac{1}{1-0.5} = 1$$

Probability Matrix Representation



P:

	s0	s1	s2	s3
s0	0	1	0	0
s1	0	0.01	0.01	0.98
s2	1	0	0	0
s3	0	0	0	1

$P_{i,k}$ = the value at the i -th row and k -th column. It specifies the probability of taking the transition $s_i \rightarrow s_k$, if we are now at s_i .

For example the circle red value above is $P_{1,2}$, specifying the probability of taking the transition from s_1 to s_2 (check the picture), which is 0.01.

Basic Operations on Probability Matrix

- Multiplying P with itself: P^n
- Multiplying a vector with P : $u \times P$
- Multiplying P with a vector: $P \times v$

Pⁿ

- $P^0 = I$ (identity matrix)
 $P^{n+1} = P \times P^n$
- $P^n_{i,k}$ is the probability of ending up in state s_k in n -steps, given we start in the state s_i .
- For example, with the previous P , let's look at P^2 :


0	1	0	0		0	1	0	0		?	?	0.01	?
0	0.01	0.01	0.98		0	0.01	0.01	0.98		?	?	?	?
1	0	0	0		1	0	0	0		?	?	?	?
0	0	0	1		0	0	0	1		?	?	?	?

$$P^2_{0,2} = (P \times P)_{0,2}$$

$$= P_{0,0} * P_{0,2} + P_{0,1} * P_{1,2} + P_{0,2} * P_{2,2} + P_{0,3} * P_{3,2}$$

Probability distribution of the next state, given the current distribution

- A **probability distribution** of the current state is the probability of currently being in various states. It can be given by a vector of size K , if K is the number of possible states. E.g. if $u = [0, 0.5, 0.5, 0]$ is the probability distribution of the current state, it says e.g. that there is 0.5 probability that currently we are in the state s_1 , but 0 probability that we are in the state s_0 .
- The product $u \times P$ (we often simply write it as uP) gives a new vector u' of size K , that gives us the probability distribution of the next state.



e.g. $u'_1 = u \cdot \text{the green column (dot product)}$

$$= u_0 * P_{0,1} + u_1 * P_{1,1} + u_2 * P_{2,1} + u_3 * P_{3,1}$$

Probability vector

- Sometimes we also want to know what the probability to end up in state, say, s_1 or s_2 as the **next** state, if we start in the state s_1 .
- We can represent “end up in either s_1 or s_2 ” with a vector $v = [0, 1, 1, 0]$.
- Let v^t is the *transpose* of v . The product $P \times v^t$ gives a w such that w is a (transposed) vector, where w_i is the probability to end up in one of the states specified in v , if we start in s_i .

0	1	0	0
0	0.01	0.01	0.98
1	0	0	0
0	0	0	1

 $\times$

0
1
1
0



1
0.02
0
0

v^t w

e.g. w_1 = the green row $\cdot v^t$ (dot product)
 $= P_{1,0} * v_0 + P_{1,1} * v_1 + P_{1,2} * v_2 + P_{1,3} * v_3$

Probability vector


- Prob(φ) (notice the underscore) is a probability vector e.g. $w =$

1
0.02
0
0

such that the i -th element tells us what the probability that the system would behave as φ if executed in state s_i .

- Example: the above w (blue) happens to be equal to Prob($\mathbf{X}(\text{try } v \text{ fail})$).
- This notation Prob will be used later when we discuss model checking of probabilistic-CTL.

The construction of P' for bounded Until




0	1	0	0
0	0.01	0.01	0.98
1	0	0	0
0	0	0	1

The probability matrix P of the DTMC on the left.

- Consider as an example to check whether the DTMC satisfies $P_{>0.99}[\text{try} \vee \neg\text{fail} \ U^{\leq 2} \text{succ}]$.
 - Calculate first the probability vector **Prob** $[\text{try} \vee \neg\text{fail} \ U^{\leq 2} \text{succ}]$.
 - From there you can calculate the set **Sat** $(P_{>0.99}[\text{try} \vee \neg\text{fail} \ U^{\leq 2} \text{succ}])$.
 - If the initial state s_0 is in the blue Sat-set then the property $P_{>0.99}[\text{try} \vee \neg\text{fail} \ U^{\leq 2} \text{succ}]$ holds on the DTMC.**

The construction of P' for bounded Until




0	1	0	0
0	0.01	0.01	0.98
1	0	0	0
0	0	0	1

The probability matrix P of the DTMC on the left.

- To calculate the probability vector $\text{Prob}[\text{try} \vee \neg\text{fail} \ U^{\leq 2} \text{succ}]$, we would like to use the matrix P above, however it will also "contain" transitions that cause you to break the green-property. So the idea is to use a "modified" matrix P' .
- We pre-calculate first the $S^{\text{yes}} = \text{Sat}(\text{succ}) = \{s3\}$. On all states in S^{yes} , you have the green property immediately (in 0 step).
- We pre-calculate S^{no} , we take $S^{\text{no}} = \text{Sat}(\neg(\text{try} \vee \neg\text{fail}) \wedge \neg \text{succ}) = \{s2\}$. Executions starting from S^{no} won't satisfy your green-property above,

The construction of P' for bounded Until



$P :$


0	1	0	0
0	0.01	0.01	0.98
1	0	0	0
0	0	0	1

$P' :$

0	1	0	0
0	0.01	0.01	0.98
0	0	0	0
0	0	0	1

1. We remove outgoing arrows from the states in S^{no} and S^{yes} .
2. We keep all arrows that go out from states which are **not** in S^{no} nor S^{yes} .
3. We add a self-loop $s \rightarrow s$ with probability 1 for any state s in S^{yes} .

Using P' for bounded Until



P' :

0	1	0	0
0	0.01	0.01	0.98
0	0	0	0
0	0	0	1

- We now use P' to iteratively calculate $\text{Prob}[\text{try } \vee \neg\text{fail } \mathbf{U}^{\leq 2} \text{succ}]$

- From S^{yes} you know that $\text{Prob}[\text{try } \vee \neg\text{fail } \mathbf{U}^{\leq 0} \text{succ}] =$

0
0
0
1

- $\text{Prob}[\text{try } \vee \neg\text{fail } \mathbf{U}^{\leq 1} \text{succ}] = P' \times$

0
0
0
1

 $=$

0
0.98
0
1


- $\text{Prob}[\text{try } \vee \neg\text{fail } \mathbf{U}^{\leq 2} \text{succ}] = P' \times$

0
0.98
0
1

 $=$

0.98
0.9898
0
1


So, does the property hold?



- We have calculated $\text{Prob}[\text{try} \vee \neg\text{fail} \text{ } \mathbf{U}^{\leq 2} \text{ succ}] = [0.98, 0.9898, 0, 1]$
- So, the set $\text{Sat}(P_{>0.99}[\text{try} \vee \neg\text{fail} \text{ } \mathbf{U}^{\leq 2} \text{ succ}]) = \{s_3\}$
- So we conclude that the DTMC does **not** satisfy the claimed property $P_{>0.99}[\text{try} \vee \neg\text{fail} \text{ } \mathbf{U}^{\leq 2} \text{ succ}]$.


Buchi vs Rabin

- Consider the LTL property $\Diamond \Box a$. This can be described by this Buchi automaton, with $\{1\}$ as the accepting state:



Notice that this Buchi is **non-deterministic**. As such, we can use it for model checking on a probabilistic model such as DTMC.


- We can however represent the property with a **deterministic** Rabin automaton, with the pair $(\{0\}, \{1\})$ as its accepting condition.




- Note that every deterministic Buchi automaton is also a Rabin automaton.

Example

Consider the DTMC M below. We want to check if the property $P_{>0.99}[\varphi]$ holds on M , where LTL $\varphi = \mathbf{X}(\text{try } \mathbf{W} \text{ succ})$




First, we construct a deterministic Rabin automaton that equivalently describes φ . I give a deterministic Buchi. For each state, you also need to cover **all possible** outgoing arrows.



As a Buchi, it is also a Rabin, with accepting condition: $\{(L, K)\}$ with $L = \emptyset$ and $K = \{v, w\}$. The state x is a "sink" state; ending up there will deny establishing φ . The state x is added to make outgoing arrows from v complete (defined for all possible combination of p and q).


Intersection DTMC and Rabin

Before we construct the intersection $M \cap R$ (or product $M \otimes R$, in Kwiatkowska's notation), I prefer to first transform M so that we have labels on the arrows, by moving them to incoming arrows. The probabilities of the transitions remain the same, of course.



Intersection DTMC and Rabin


Now we can construct the intersection. This is done as we did with standard LTL, with as addition that probabilities are kept as in M:



This results in a new Rabin R2. Any state (s, t) in R2 with t in R's L is marked as L in R2. If t was in R's K, we mark (s, t) as K in R2. So, for above example we get L is empty, and all the blue states are K.

Calculating $\text{Prob}(\varphi)$

So we have this $R2 = M \cap R$, with $L = \emptyset$ and $K = \{ \text{the blue states} \}$. Any execution satisfying φ can only pass L finitely many often and should pass one of K infinitely many often.



The colored boxes are the **bottom** SCCs of $R2$. Only the blue one, let's call it Z , contains a member of K . And: $Z \cap L = \emptyset$. So, any execution satisfying φ will end up in Z .

We know:

- $\text{Prob}((s_3, w), \varphi) = 1$. For the red-states: $\text{Prob}((\dots, x), \varphi) = 0$
- $\text{Prob}((z, u), \varphi) = \text{Prob}((s_0, v), \varphi) = \text{Prob}((s_1, v), \varphi)$
- $\text{Prob}((s_1, v), \varphi) = 0.01 * \text{Prob}((s_1, v), \varphi) + 0.98 * \text{Prob}((s_3, w), \varphi)$

$= 1$

Solving the equations gives $\text{Prob}((z, u), \varphi) = \text{Prob}((s_1, v), \varphi) = 98/99 < 0.99$.

So, the answer to the original question if $P_{>0.99}[\varphi]$ holds is a “no”.

additional slides MDP

Calculating $\text{pmin}(s, \varphi_1 \mathbf{U} \varphi_2)$

$\text{pmin}(s, \varphi)$ = the **minimum** probability for having executions starting from s satisfying φ , regardless the adversary.

To calculate $\text{pmin}(s, \varphi_1 \mathbf{U} \varphi_2)$:

1. Calculate first the sets S^{yes} and S^{no} .
 - $S^{\text{yes}} = \{ s \mid P(s, \varphi_1 \mathbf{U} \varphi_2) \geq 1, \text{ for all adversaries } \} \rightarrow \text{with algorithm Prob1A.}$
 - $S^{\text{no}} = \{ s \mid P(s, \varphi_1 \mathbf{U} \varphi_2) \leq 0, \text{ for some adversary } \} \rightarrow \text{with algorithm Prob0E.}$
2. For any state s in S^{yes} , we then know that $\text{pmin}(s, \varphi_1 \mathbf{U} \varphi_2) \geq 1$.
3. For any state s in S^{no} we have $\text{pmin}(s, \varphi_1 \mathbf{U} \varphi_2) \leq 0$.
4. We then proceed with calculating the pmin for the remaining states (which are not in S^{yes} nor S^{no}).

Algorithm Prob0E

- The algorithm below first calculates the set R of all states s satisfying $E[s, \varphi_1 \cup \varphi_2]$, regardless the adversary. So, for any state in R , and for any adversary $\text{Prob}(s, \varphi_1 \cup \varphi_2) > 0$.
- S^{no} is just complement S/R .
- $\text{Sat}(\varphi_1)$ and $\text{Sat}(\varphi_2)$ in the paremeters are the set of states on which φ_1 and φ_2 respectively hold.

$\text{PROB0E}(\text{Sat}(\phi_1), \text{Sat}(\phi_2))$

```
1.   $R := \text{Sat}(\phi_2)$ 
2.   $done := \text{false}$ 
3.  while ( $done = \text{false}$ )
4.       $R' := R \cup \{s \in \text{Sat}(\phi_1) \mid \forall \mu \in \text{Steps}(s). \exists s' \in R. \mu(s') > 0\}$ 
5.      if ( $R' = R$ ) then  $done := \text{true}$ 
6.       $R := R'$ 
7.  endwhile
8.  return  $S \setminus R$ 
```

Prob1A

- Calculate first the set F of states from where we have a path passing exclusively through states satisfying $\varphi_1 \wedge \neg\varphi_2$ and ends in S^{no} , under **some** adversary.
By definition this F also includes S^{no} .
- So any state s in F has $\text{Prob}(s, \varphi_1 \mathbf{U} \varphi_2) < 1$, for some adversary. In other words $\text{pmin}(s, \varphi_1 \mathbf{U} \varphi_2) < 1$.
- P^{yes} in the complement S/F .
- **Note:** for the calculation of $\text{pmin}(s, \varphi_1 \mathbf{U} \varphi_2)$, we can also just take $S^{\text{yes}} = \text{Sat}(\varphi_2)$. The calculation would still work, though it would take more steps to get its final results.

Calculating $\text{pmax}(s, \varphi_1 \mathbf{U} \varphi_2)$

$\text{pmax}(s, \varphi)$ = the **maximum** probability for having executions starting from s satisfying φ , regardless the adversary.

To calculate $\text{pmax}(s, \varphi_1 \mathbf{U} \varphi_2)$:

1. Calculate first the sets S^{yes} and S^{no} .
 - $S^{\text{yes}} = \{ s \mid P(s, \varphi_1 \mathbf{U} \varphi_2) \geq 1, \text{ for } \mathbf{some} \text{ adversaries} \} \rightarrow \text{with algorithm prob1E.}$
 - $S^{\text{no}} = \{ s \mid P(s, \varphi_1 \mathbf{U} \varphi_2) \leq 0, \text{ for } \mathbf{all} \text{ adversary} \} \rightarrow \text{with algorithm prob0A.}$
2. For any state s in S^{yes} , we then know that $\text{pmax}(s, \varphi_1 \mathbf{U} \varphi_2) \geq 1$.
3. For any state s in S^{no} we have $\text{pmax}(s, \varphi_1 \mathbf{U} \varphi_2) \leq 0$.
4. We then proceed with calculating the pmin for the remaining states (which are not in S^{yes} nor S^{no}).

Algorithm Prob0A

- The algorithm below first calculates the set R of all states s satisfying $E[s, \varphi_1 \cup \varphi_2]$, for some adversary. So, for any state in R , there is an adversary such that $\text{Prob}(s, \varphi_1 \cup \varphi_2) > 0$.
- S^{no} is just complement S/R .
- $\text{Sat}(\varphi_1)$ and $\text{Sat}(\varphi_2)$ in the paremeters are the set of states on which φ_1 and φ_2 respectively hold.

$\text{PROB0A}(\text{Sat}(\phi_1), \text{Sat}(\phi_2))$

```
1.   $R := \text{Sat}(\phi_2)$ 
2.   $done := \text{false}$ 
3.  while ( $done = \text{false}$ )
4.       $R' := R \cup \{s \in \text{Sat}(\phi_1) \mid \exists \mu \in \text{Steps}(s). \exists s' \in R. \mu(s') > 0\}$ 
5.      if ( $R' = R$ ) then  $done := \text{true}$ 
6.       $R := R'$ 
7.  endwhile
8.  return  $S \setminus R$ 
```

Prob1E

- More complicated. See Dave's slides on MDP.

Prob1E