

Model Checking with SPIN

A Bit More about SPIN

Wishnu Prasetya

wishnu@cs.uu.nl

www.cs.uu.nl/docs/vakken/pv

Content

- SPIN internal data structures
- SPIN result report
- Writing LTL formulas
- Containing state explosion
- Example

Acknowledgement: some slides are taken and adapted from Theo Ruys's SPIN Tutorials.

Data structures involved in SPIN DFS

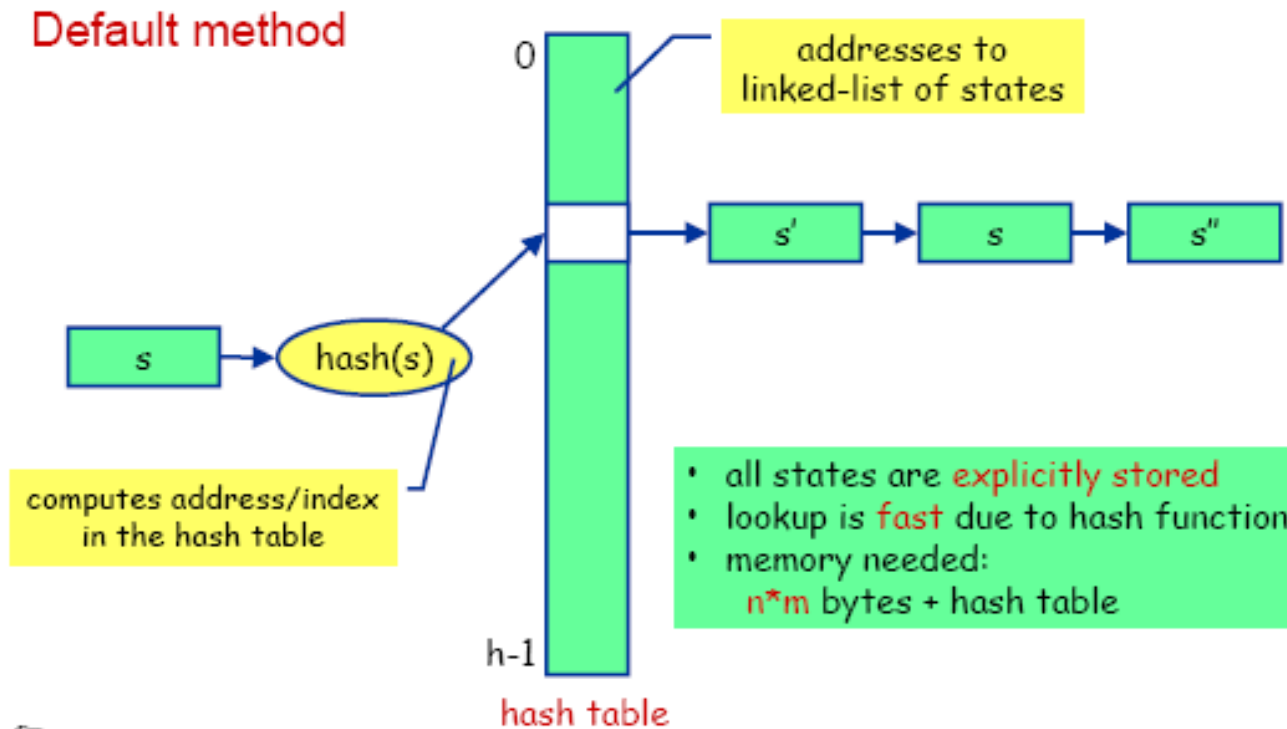
- Representation of a state.
- Stack for the DFS
 - To remember where to backtrack in DFS
 - It corresponds to the current “execution prefix” that is being inspected → used for reporting.
- Something to hold the set of visited states = “state space”.

State

- Each (global) state of a system is a tuple of the states of its processes.
- E.g. Suppose we have:
 - One global var byte x
 - Process P with byte y
 - Process Q with byte z
- A system state should describe:
 - all these variables and their values
 - Program counter of each process
 - Other SPIN predefined vars
- Such a global state can be quite big.

State-space is stored with a hash table

The list of “visited states” is maintained by a Hash-table. So matching if a state occurring in the table is fast!



Verifier's output

assertion violated !((crit[0]&&crit[1])) (at depth 5) // computation depth

...

Warning: Search not completed

Full statespace search for:

...

never-claim - (not selected)

assertion violations +

invalid endstates +

State-vector 20 byte, depth reached 7, errors: 1

24 states, stored

17 states, matched

41 transitions (= stored+matched)

// max. stack depth

// states stored in hash table

// states found re-revisited

hash conflicts: 0 (resolved)

(max size 2^{19} states)

2.542 memory usage (Mbyte)

Watch out for state explosion!

```
int x,y,z ;
```

```
P { do :: x++ od }
```

```
Q { do :: y++ od }
```

```
R { do :: x/=y → z++ od }
```

- Size of each state: > 12 bytes
- Number of possible states $\approx (2^{32})^3 = 2^{96}$
- Using byte (instead of int) brings this down to 50 MB
- Focus on the critical aspect of your model (e.g. its concurrency); abstract from data when possible.

Another source of explosion : concurrency

imposing a coarser grain atomicity

```
atomic { guard → stmt_1; ... ; stmt_n }
```

```
active proctype P { int x ; (y>0) ; y-- ; x=y }
```

└──────────────────┘
put in *atomic* ?

- more abstract, less error prone, but less parallelism
- executable if the guard statement is executable
- none of stmt-i should be blocking; or rather : if any of them blocks, atomicity is lost

d_step sequences

```
d_step { guard → stmt_1; ... ; stmt_n }
```

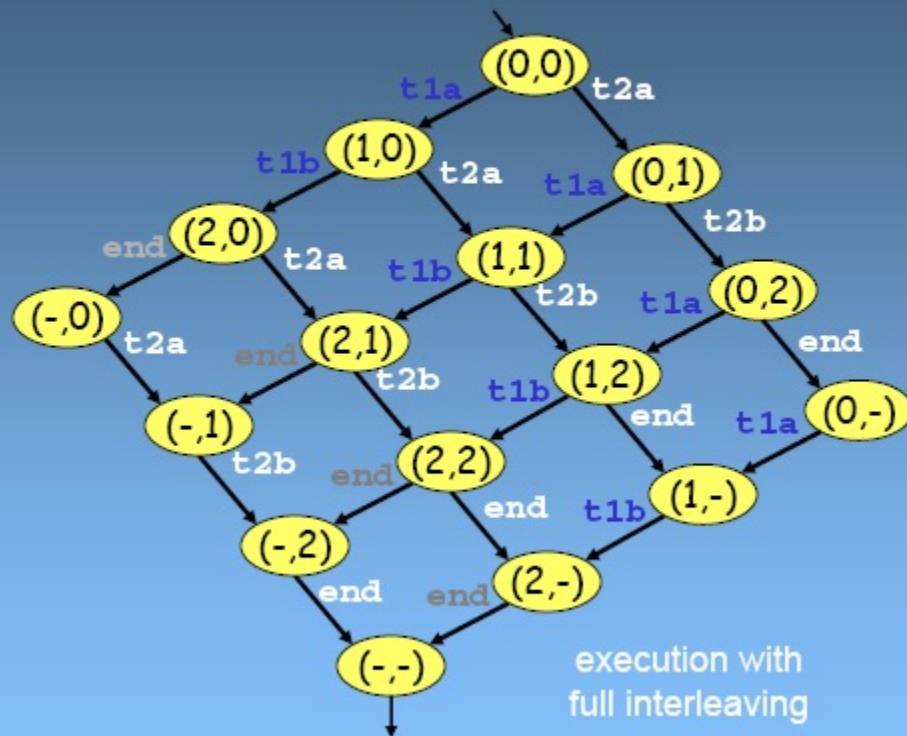
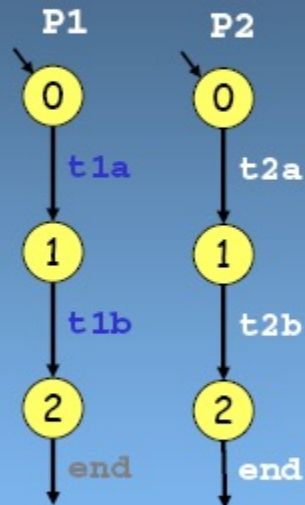
```
d_step {  /* reset array elements to 0 */  
    i = 0;  
    do  
        :: i < N -> x[i] = 0; i++  
        :: else -> break  
    od;  
    i = 0  
}
```

- like an atomic, but *must be deterministic* and *may not block anywhere*
- atomic and d_step sequences are often used as a model reduction method, to *lower complexity* of large models (improving tractability)
- No jump into the middle of a d_step

execution without atomics or d_steps

```
active proctype P1() { t1a; t1b }  
active proctype P2() { t2a; t2b }
```

execution
without atomics or d_steps



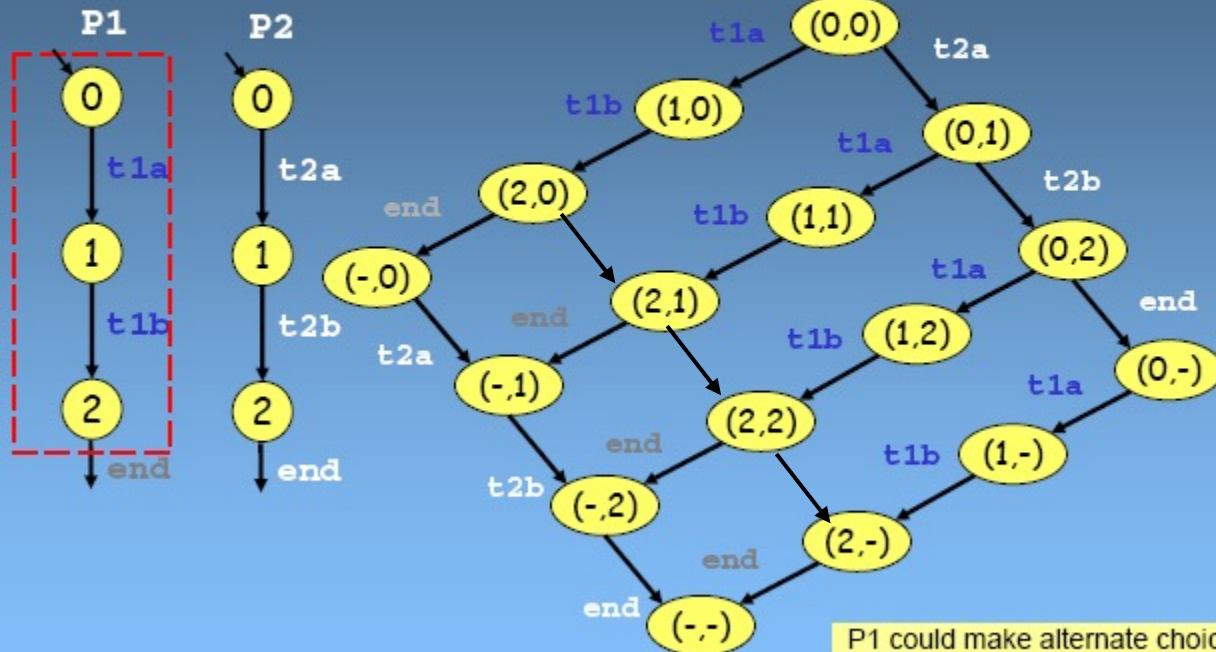
execution with
full interleaving

execution with one atomic sequence

```
active proctype P1() { atomic { t1a; t1b } }  
active proctype P2() { t2a; t2b }
```

execution with one
atomic sequence

P2 can be interrupted, but not P1



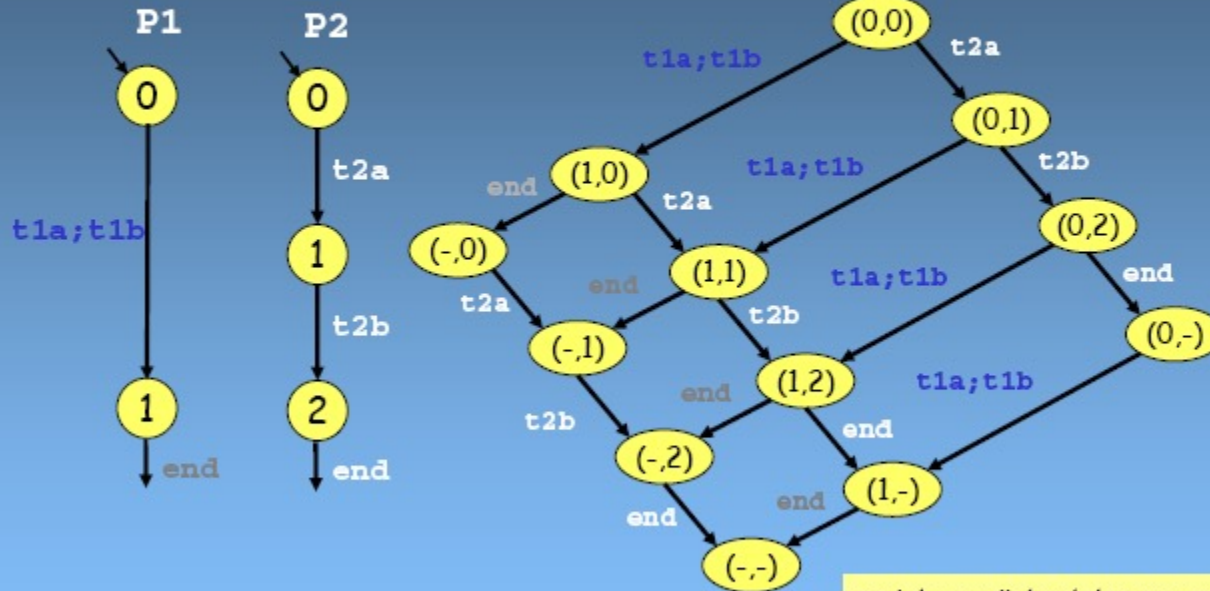
P1 could make alternate choices at the intermediate states (e.g., in if or do-statements)

execution with a d_step sequence

```
active proctype P1() { d_step {t1a; t1b} }  
active proctype P2() { t2a; t2b }
```

execution with a
d_step sequence

P1 now has only one transition...



no intermediate states are created:
faster, smaller graph, but no non-
determinism possible inside d_step
sequence itself

atomic vs d_step

- **d_step:**
 - executed as *one* block
 - deterministic
 - blocking or non-termination would hang you 😊
- **atomic:**
 - translated to a series of actions
 - executed step-by-step, but without interleaving
 - it can make non-deterministic choices

Partial Order Reduction

- The validity of a property ϕ is often insensitive to the order in which 'independent' actions are interleaved.

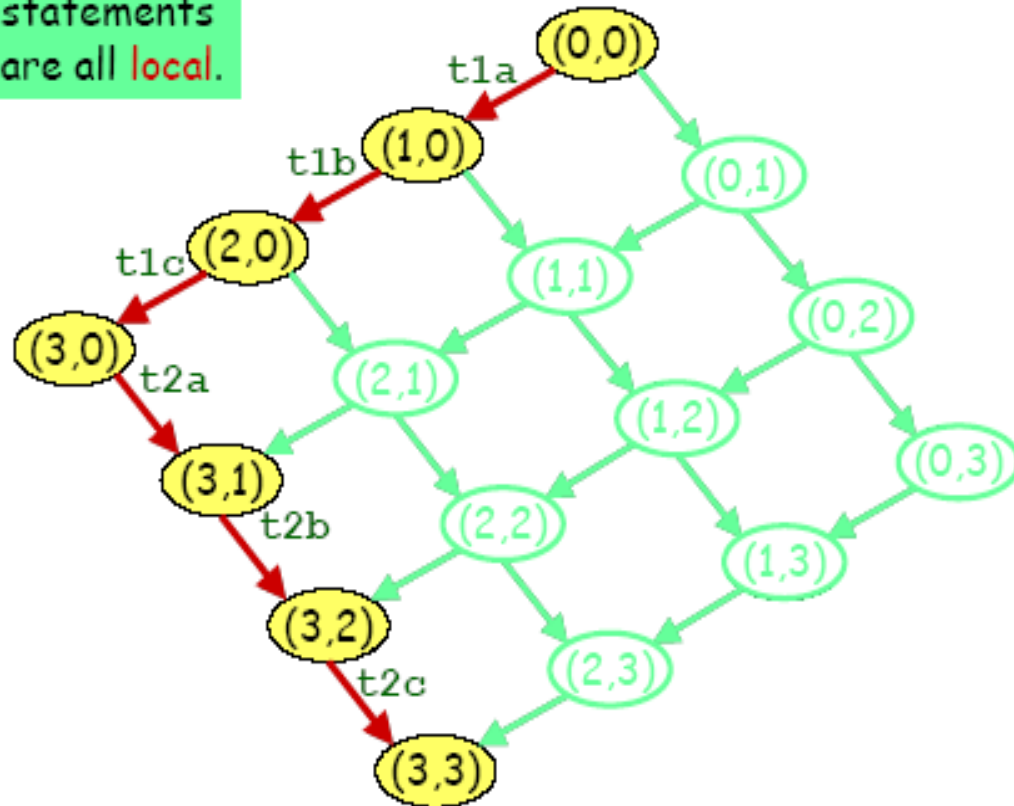
e.g. stutter invariant ϕ (does not contain X) that only refers to global variables, is insensitive to the relative order of actions that only access local variables.

- Idea: if in some global state, a process P can execute only actions updating local variables, always do these actions first (so they will not be interleaved!)
- We can also do the same with receive/send actions :
 - $c ? x$ where x is local, and this process is the only receiver of c
 - $d ! y$ where this process is the only sender to d
 - assuming ϕ does not refer to channels

Reduction Algorithms

- Partial Order Reduction

Suppose the statements of P1 and P2 are all **local**.



Results on Partial Order Reduction

Protocol	Algorithm	States	Transitions	Time(sec.)	Memory (Mb)
Best-Case	Non-Reduced	100,001	450,002	13.2	4.3
	Static Reduction	47	47	(<0.1)	1.0
	Dynamic Reduction	47	47	0.1	1.4
Worst-Case	Non-Reduced	100,001	450,002	14.5	5.0
	Static Reduction	100,001	450,002	16.7	5.1
	Dynamic Reduction	100,001	450,002	84.5	5.3
Tpc	Non-Reduced	3,918,286	11,762,426	630.6	268.4
	Static Reduction	391,534	466,753	30.6	26.2
	Dynamic Reduction	267,204	295,395	131.4	18.9
Snoopy	Non-Reduced	91,920	305,460	14.4	11.5
	Static Reduction	16,279	23,532	1.7	3.2
	Dynamic Reduction	7,158	8,459	6.8	2.6
Pftp	Non-Reduced	417,321	1,244,865	73.2	62.3
	Static Reduction	53,244	67,901	6.8	9.3
	Dynamic Reduction	125,718	163,459	105.5	20.6
Leader	Non-Reduced	45,885	185,032	8.1	9.6
	Static Reduction	79	79	0.1	1.1
	Dynamic Reduction	79	79	0.2	1.4

This result is from Holzmann & Peled in 1994, on a Sparc-10 workstation with 128Mbyte of RAM. (about 40 mhz; so 1 mips??)

Specifying LTL properties

- (Check out the Manual)

```
#define PinCritical    crit[1]
#define QinCritical    crit[2]

ltl name { []!(PinCritical && QinCritical) }
```

- SPIN then generates the Buchi automaton for this LTL formula; called “Never Claim” in SPIN.

Example of a Never Claim

To verify: $\langle \rangle [] p$

SPIN generates this never-claim / Buchi of $[] \langle \rangle \neg p$

```
never {  
  init:  
    if  
    ::  $\neg p \rightarrow$  goto accept  
    :: else  $\rightarrow$  goto init  
    fi;  
  accept:  
    skip; goto init ;  
}
```

Neverclaim

- From SPIN perspective, a neverclaim (NC) is just another process, but it is executed in “lock-step”. See manual; essentially it means :
 - Each time the system takes a step, the NC also takes a step (but only an enabled step).
 - If one of them cannot make a step, the lock-step execution stops.
- You can use it to “match” an execution
 - To match an infinite execution, as in Buchi automaton
 - The NC reaches its final state (its final “}”) → match! → used to match against finite executions.
 - We used it to find a counter example, but you can use it as a way to query your state space.

You can also manually write a custom NC to query the set of possible executions of your model

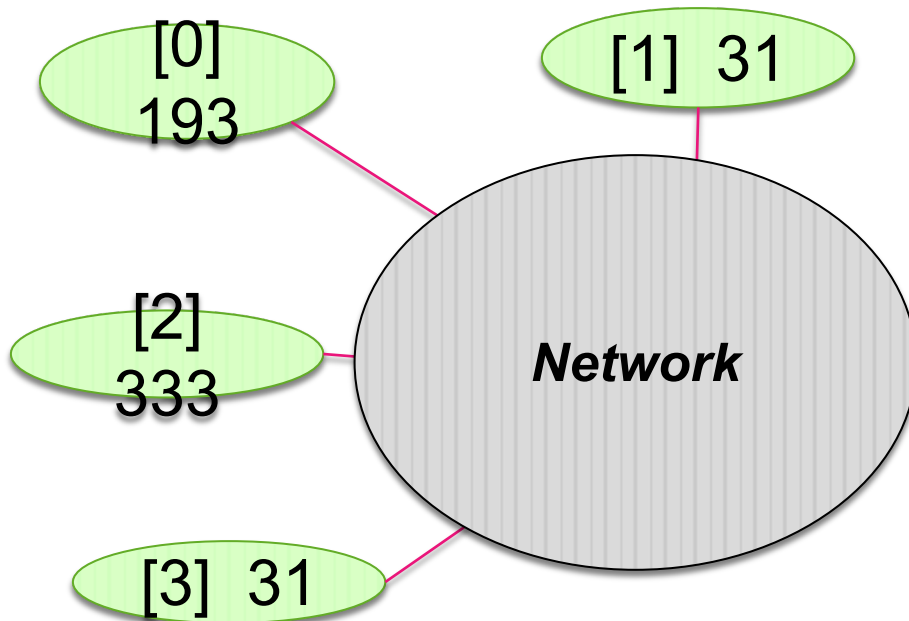
```
never {  
  do  
    :: (b) ; (!b) ; break  
    :: skip  
  od  
  accept:  
  skip ; goto accept ;  
}
```

Will match with an execution where at some point b holds, then at the next step it does not hold.

```
never {  
  accept : do :: (x==0) ; (x==1) od  
}
```

Match with an execution where $(x==0)(x==1)$ holds alternatingly,

Example: distributed sorting



- Idea:

Let $P(i)$ swap values with $P(i+1)$, if they are in the wrong order.

- Spec:

Eventually the values will be sorted.

SPIN model

```
#define N 5
```

```
byte a[N] ;
```

```
proctype P(byte i) {
```

```
    byte tmp = 0 ;
```

```
    do
```

```
    :: d_step{ a[i]>a[i+1] ->
```

```
        tmp=a[i] ;
```

```
        a[i]=a[i+1] ;
```

```
        a[i+1]=tmp ;
```

```
        tmp=0 }
```

```
    od ;
```

```
}
```

```
init {
```

```
    byte i ;
```

```
    do
```

```
    :: i<N ->
```

```
        if
```

```
        :: a[i]=0
```

```
        :: a[i]=1
```

```
        ...
```

```
        fi ; i++
```

```
    :: else -> break ;
```

```
    od ;
```

```
    i=0 ;
```

```
    do
```

```
    :: i< N - 1 -> run P(i) ; i++
```

```
    :: else -> run detect() ; break
```

```
    od
```

```
}
```

(let's just assume locking $a[i]$ and $a[i+1]$ atomically is reasonable.)

Swap values, set tmp back to 0 to save state.

Expressing the spec

Eventually the values will be sorted.

With LTL: $\langle \rangle [] (\forall i : 0 \leq i < N-1 : a[i] \leq a[i+1])$

Note: SPIN does **not** however support quantification in its Expr!

You can still manually write a neverclaim, where you implement the quantification (note that the quantification is meant to be evaluated atomically).

Detecting “termination”

New spec: we want the processes themselves to know, or to be informed, that the goal (to sort values) has been accomplished.

```
proctype detect() {  
  byte i ;  
  timeout ->  
    do  
    :: i < N-1 -> assert (a[i] <= a[i+1])  
    :: else -> break  
  od  
}
```

(Central server solution; not a distributed solution)

done = true

Extend $P(i)$, such that when it sees “done” is true, it will terminate.

Unfortunately, not good enough. The above solution uses “timeout” which in SPIN is actually implemented as a mechanism to detect non-progress; in the actual system we now assume not to have this mechanism in the first place, and hence have to implement it ourselves.

Detecting “termination”

Idea: let “detect” keep scanning the array to check if it is sorted.

```
proctype detect() {  
  byte i ;  
  i=0 ;  
  do  
  :: i<N-1 -> if  
    :: a[i]>a[i+1] -> i=0  
    :: else          -> i++  
    fi  
  :: else -> done=true ; break  
  od  
}
```

(Central server solution; not a distributed solution)

Unfortunately, this doesn't work perfectly. Consider this sequence of steps:

[4, 5, 1]

detect 0,1 → ok

swap 1,2

[4, 1, 5]

detect 1,2 → ok

now “detect” concludes termination!

Can you find a solution for this??